

# Web of Things (WoT) Architecture

## 1.1



W3C Editor's Draft 30 November 2022

### ▼ More details about this document

#### This version:

<https://w3c.github.io/wot-architecture/>

#### Latest published version:

<https://www.w3.org/TR/wot-architecture11/>

#### Latest editor's draft:

<https://w3c.github.io/wot-architecture/>

#### History:

<https://www.w3.org/standards/history/wot-architecture11>

[Commit history](#)

#### Implementation report:

<https://w3c.github.io/wot-architecture/testing/report11.html>

#### Editors:

Michael Lagally ([Oracle Corp.](#))

Ryuichi Matsukura ([Fujitsu Ltd.](#))

Michael McCool ([Intel Corp.](#))

Kunihiko Toumura ([Hitachi, Ltd.](#))

#### Former editors:

Kazuo Kajimoto (when at Panasonic)

Toru Kawaguchi ([Panasonic Corp.](#))

Matthias Kovatsch ([Huawei](#))

#### Feedback:

[GitHub w3c/wot-architecture](#) ([pull requests](#), [new issue](#), [open issues](#))

[public-wot-wg@w3.org](mailto:public-wot-wg@w3.org) with subject line [wot-architecture11] ... *message topic* ...  
([archives](#))

#### Contributors

[In the GitHub repository](#)

Copyright © 2017-2022 W3C® ([MIT](#), [ERCIM](#), [Keio](#), [Beihang](#)). W3C [liability](#), [trademark](#) and [permissive document license](#) rules apply.

---

## Abstract

The ~~W3C~~ Web of Things (WoT) enables interoperability across IoT platforms and application domains. The goal of the WoT is to preserve and complement existing IoT standards and solutions. The ~~W3C~~ WoT architecture is designed to describe what exists, and only prescribes new mechanisms when necessary.

This *WoT Architecture* specification describes the abstract architecture for the ~~W3C~~ Web of Things. This abstract architecture is based on requirements that were derived from use cases for multiple application domains. Several modular building blocks were identified whose detailed specifications are given in other documents. This document describes how these building blocks are related and work together. The WoT abstract architecture defines a basic conceptual framework that can be mapped onto a variety of concrete deployment scenarios, several examples of which are given. However, the abstract architecture described in this specification does not itself define concrete mechanisms or prescribe any concrete implementation.

- Sections 4 and 5 ([Application Domains](#) and [Common Deployment Patterns](#)) explain ways to use WoT. These sections are self contained, do not contain any assertion and other specifications do not depend on these sections.
- Section 6 [Abstract WoT System Architecture](#), explains architectural elements of the ~~W3C~~ WoT. This section is normative and contains assertions that are relevant for WoT implementations.
- Section 7 [WoT Building Blocks](#), explains building blocks of the ~~W3C~~ WoT.
- Section 8 [Abstract Servient Architecture](#) describes the architecture of a [Servient](#). It informatively describes how to implement a [WoT Runtime](#) on a device.
- Section 9 [Example WoT Deployments](#) provides various examples of how the Web of Things (WoT) abstract architecture may be instantiated when devices and services that implement the Thing and Consumer roles interact in different network topologies.
- Section 10 [Security Considerations](#) and Section 11 [Privacy Considerations](#) are normative; they summarize some general issues and provide guidelines to help preserve the security and privacy of concrete WoT implementations.

## Status of This Document

*This section describes the status of this document at the time of its publication. A list of current ~~W3C~~ publications and the latest revision of this technical report can be found in the [W3C technical reports index](#) at <https://www.w3.org/TR/>.*

This document describes an abstract architecture. However, there is an [Implementation Report](#) that describes a set of concrete implementations following the ~~W3C~~ Web of Things architecture. It also references the other implementation reports for the various WoT building blocks.

The Web of Things Working Group intends to submit this document for consideration as a W3C Proposed Recommendation after at least the minimum CR review period has passed. However, before PR transition is requested, any features or assertions currently marked as at-risk that did not appear in the Architecture 1.0 specification and do not have at least two implementations at that time will either be removed or converted into informative statements, as appropriate.

At-risk assertions are marked with yellow highlighting.

This document was published by the [Web of Things Working Group](#) as an Editor's Draft.

Publication as an Editor's Draft does not imply endorsement by W3C and its Members.

This is a draft document and may be updated, replaced or obsoleted by other documents at any time. It is inappropriate to cite this document as other than work in progress.

This document was produced by a group operating under the [W3C Patent Policy](#). W3C maintains a [public list of any patent disclosures](#) made in connection with the deliverables of the group; that page also includes instructions for disclosing a patent. An individual who has actual knowledge of a patent which the individual believes contains [Essential Claim\(s\)](#) must disclose the information in accordance with [section 6 of the W3C Patent Policy](#).

This document is governed by the [2 November 2021 W3C Process Document](#).

## Table of Contents

### **Abstract**

### **Status of This Document**

#### **1. Introduction**

#### **2. Conformance**

#### **3. Terminology**

##### 3.1 Device Categories

#### **4. Application Domains**

##### 4.1 Consumer

##### 4.2 Industrial

##### 4.2.1 Example: Smart Factory

##### 4.3 Transportation & Logistics

##### 4.4 Utilities

##### 4.5 Oil and Gas

- 4.6 Insurance
- 4.7 Engineering and Construction
- 4.8 Agriculture
- 4.9 Healthcare
- 4.10 Environment Monitoring
- 4.11 Smart Cities
- 4.12 Smart Buildings
- 4.13 Connected Car
- 4.13.1 Connected Car Example

## **5. Common Deployment Patterns**

- 5.1 Telemetry
- 5.2 Device Controllers
- 5.3 Thing-to-Thing
- 5.4 Remote Access
- 5.5 Smart Home Gateways
- 5.6 Edge Devices
- 5.7 Digital Twins
  - 5.7.1 Cloud Connected Devices
  - 5.7.2 Orchestration
  - 5.7.3 Legacy Devices
- 5.8 Multi-Cloud
- 5.9 Virtual Things
- 5.10 Cross-domain Collaboration
- 5.11 System Integration

## **6. Abstract WoT System Architecture**

- 6.1 Fundamental Concepts
  - 6.1.1 Metadata
    - 6.1.1.1 Thing Descriptions
    - 6.1.1.2 Thing Models
  - 6.1.2 Links
  - 6.1.3 Intermediaries
- 6.2 Affordances
- 6.3 Web Thing
- 6.4 Lifecycle
  - 6.4.1 Simple System Lifecycle
  - 6.4.2 System Lifecycle with Registration
  - 6.4.3 Thing Lifecycle
- 6.5 Interaction Model
  - 6.5.1 Properties

6.5.2	Actions
6.5.3	Events
6.6	Hypermedia Controls
6.6.1	Links
6.6.2	Forms
6.7	Protocol Bindings
6.7.1	Hypermedia-driven
6.7.2	URIs
6.8	Media Types
6.9	Internationalization
6.10	WoT System Components and their Interconnectivity
6.10.1	Direct Communication
6.10.2	Indirect Communication
<b>7.</b>	<b>WoT Building Blocks</b>
7.1	WoT Thing Description
7.2	Thing Model
7.3	Profiles
7.3.1	Profiling Methodology
7.4	WoT Discovery
7.4.1	Introduction Mechanisms
7.4.2	Exploration Mechanisms
7.5	WoT Binding Templates
7.6	WoT Scripting API
7.7	WoT Security and Privacy Guidelines
<b>8.</b>	<b>Abstract Servient Architecture</b>
8.1	Behavior Implementation
8.2	WoT Runtime
8.3	WoT Scripting API
8.4	Exposed Thing and Consumed Thing Abstractions
8.5	Private Security Data
8.6	Protocol Stack Implementation
8.7	System API
8.8	Alternative Servient and WoT Implementations
8.8.1	Native WoT API
8.8.2	Thing Description for Existing Devices
<b>9.</b>	<b>Example WoT Deployments</b>
9.1	Thing and Consumer Roles
9.2	Topology of WoT Systems and Deployment Scenarios

- 9.2.1 Consumer and Thing on the Same Network
- 9.2.2 Consumer and Thing Connected via Intermediaries
  - 9.2.2.1 Intermediary Acting as a Proxy
  - 9.2.2.2 Intermediary Acting as a Digital Twin
- 9.2.3 Devices in a Local Network Controlled from a Cloud Service
- 9.2.4 Discovery Using a Thing Description Directory
- 9.2.5 Service-to-Service Connections Across Multiple Domains
  - 9.2.5.1 Connection Through Thing Description Directory Synchronization
  - 9.2.5.2 Connection Through Proxy Synchronization

## **10. Security Considerations**

- 10.1 WoT Thing Description Risks
  - 10.1.1 Thing Description Private Security Data Risk
  - 10.1.2 Thing Description Communication Metadata Risk
- 10.2 WoT Scripting API Risks
  - 10.2.1 Cross-Script Security Risk
  - 10.2.2 Physical Device Direct Access Risk
- 10.3 WoT Runtime Risks
  - 10.3.1 Provisioning and Update Security Risk
  - 10.3.2 Security Credentials Storage Risk
- 10.4 Trusted Environment Risks
- 10.5 Secure Transport

## **11. Privacy Considerations**

- 11.1 WoT Thing Description Risks
  - 11.1.1 Thing Description Personally Identifiable Information Risk
- 11.2 Access to Personally Identifiable Information

## **A. Recent Specification Changes**

- A.1 Changes from the WD published September 7, 2022
- A.2 Changes in WD published September 7, 2022 from the FPWD version
- A.3 Changes in the FPWD from the 1.0 version of [*wot-architecture*]

## **B. Acknowledgments**

## **C. References**

- C.1 Normative references
- C.2 Informative references

*This section is non-normative.*

The goals of the *Web of Things* (WoT) are to improve the interoperability and usability of the Internet of Things (IoT). Through a collaboration involving many stakeholders over many years, several building blocks have been identified that help address these challenges.

A set of over 30 WoT *use cases* were contributed by stakeholders from multiple industries for various application domains. These have been collected and were published in the *WoT Use Cases and Requirements* <https://www.w3.org/TR/wot-usecases/> document.

The collection of use cases is classified into two categories:

- Horizontal use cases that address multiple domains
- Domain specific (vertical) use cases for a single application domain

These use cases and requirements drive the creation and further evolution of the W3C WoT specification family.

The WoT architecture specification is focused on the scope of W3C WoT standardization, which can be broken down into these building blocks as well as the abstract architecture that defines how they are related.

The architecture document serves multiple purposes:

The building blocks are defined and described in detail in separate specifications. In addition to defining the abstract architecture and its terminology and conceptual framework, this specification also serves as an introduction to the WoT building blocks, and explains their interworking:

- The *Web of Things (WoT) Thing Description* [[WOT-THING-DESCRIPTION](#)] normatively provides a machine-readable data format for describing the metadata and network-facing interfaces of Things. It is based upon the fundamental concepts introduced in this document, such as interaction affordances.
- The *Web of Things (WoT) Binding Templates* [[WOT-BINDING-TEMPLATES](#)] provides informational guidelines on how to define network-facing interfaces in Things for particular protocols and IoT ecosystems, which we call Protocol Bindings. The document also provides examples for a number of existing IoT ecosystems and standards.
- The *Web of Things (WoT) Discovery* [[WOT-DISCOVERY](#)] specification defines a distribution mechanism for WoT metadata (Thing Descriptions). The WoT Discovery process uses existing mechanisms for first contact, but provides for access control before serving detailed metadata. It includes support for directories and self-description.
- The *Web of Things (WoT) Scripting API* [[WOT-SCRIPTING-API](#)], which is optional, enables the implementation of the application logic of a Thing using a common JavaScript API similar

to the Web browser APIs. This simplifies IoT application development and enables portability across vendors and devices.

- The *Web of Things (WoT) Security and Privacy Guidelines* [WOT-SECURITY] represent a cross-cutting building block. This informational document provides guidelines for the secure implementation and configuration of Things, and discusses issues which should be considered in any systems implementing W3C WoT. However, it should be emphasized that security and privacy can only be fully evaluated in the context of a complete set of concrete mechanisms for a specific implementation, which goes beyond the scope of the WoT abstract architecture. This is especially true when the WoT architecture is used descriptively for pre-existing systems, since the W3C WoT cannot constrain the behavior of such systems, it can only describe them. In this document we also discuss privacy and security risks and their mitigation at a high level in sections 10. Security Considerations and 11. Privacy Considerations.

This specification also covers non-normative architectural aspects and conditions for the deployment of WoT systems. These guidelines are described in the context of example deployment scenarios, although this specification does not require specific concrete implementations.

This specification serves as an umbrella for W3C WoT specifications and defines the basics such as terminology and the underlying abstract architecture of the W3C Web of Things. In summary, the purpose of this specification is to provide:

- a set of application domains in 4. Application Domains that were considered to identify use cases for the W3C WoT Architecture,
- a set of common deployment patterns in 5. Common Deployment Patterns,
- a definition of the abstract architecture in 6. Abstract WoT System Architecture,
- an overview of a set of WoT building blocks and their interplay in 7. WoT Building Blocks,
- an informative guideline on how to map the abstract architecture to concrete implementations in 8. Abstract Servient Architecture,
- informative examples of deployment scenarios in 9. Example WoT Deployments,
- and a set of high level security and privacy considerations to be aware of when implementing a system based on the W3C WoT architecture in 10. Security Considerations and 11. Privacy Considerations, respectively.

Additional requirements, use cases, conceptual features and new building blocks are collected in future versions of the *WoT Use Cases and Requirement* <https://www.w3.org/TR/wot-usecases/> document.

## § 2. Conformance



As well as sections marked as non-normative, all authoring guidelines, diagrams, examples, and notes in this specification are non-normative. Everything else in this specification is normative.

The key words *MAY*, *MUST*, *MUST NOT*, *SHOULD*, and *SHOULD NOT* in this document are to be interpreted as described in [BCP 14](#) [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

## § 3. Terminology

*This section is non-normative.*

This specification uses the following terms as defined here. The WoT prefix is used to avoid ambiguity for terms that are (re)defined specifically for Web of Things concepts.

In case of a conflict of a definition with terminology used in another WoT document, the definition of the WoT Architecture takes precedence.

### **Action**

An Interaction Affordance that allows to invoke a function of the Thing, which manipulates state (e.g., toggling a lamp on or off) or triggers a process on the Thing (e.g., dim a lamp over time).

### **Anonymous TD**

A Thing Description without a user-defined identifier (`id` attribute).

### **Connected Device**

A synonym for [Device](#).

### **Binding Templates**

A re-usable collection of blueprints that enable a Thing Description to be used with a specific protocol, data payload format or an IoT platform that combine both of them in specific ways. This is done through additional descriptive vocabularies, Thing Models and examples that aim to guide the implementers of Things and Consumers alike.

### **Consumed Thing**

A software abstraction that represents a remote Thing used by the local application. The abstraction might be created by a native WoT Runtime, or instantiated as an object through the WoT Scripting API.

### **Content Type**

Identifier for the format of the message body. Also known as media type and MIME type [RFC2046].

### **Consuming a Thing**

To parse and process a TD document and from it create a Consumed Thing software abstraction as interface for the application in the local runtime environment.

**Consumer**

An entity that can process WoT Thing Descriptions (including its JSON-based representation format) and interact with Things (i.e., consume Things).

**Data Schema**

A data schema describes the information model and the related payload structure and corresponding data items that are passed between Things and Consumers during interactions.

**Device**

A Device is a physical entity that has a network interface. Devices can be described by a Thing Description and are a kind of Thing. A synonym for Connected Device. Compare with Service.

**Digital Twin**

A digital twin is type of Virtual Thing that resides on a cloud or edge node. Digital Twins may be used to represent and provide a network interface for real-world devices which may not be continuously online (see also Shadows), may be able to run simulations of new applications and services before they get deployed to the real devices, may be able to maintain a history of past state or behaviour, and may be able to predict future state or behaviour. Digital Twins typically have more functionality than simple Shadows.

**Directory**

A Service that maintains a set of data or metadata describing other Services or Things. An example would be a WoT Thing Description Directory.

**Discovery**

Mechanisms defined by WoT for distributing and accessing WoT Thing Descriptions on the network, either locally or remotely.

**Discoverer**

An entity which acts as a client of a WoT Discovery process to discover and fetch a Thing Description, e.g. by being introduced to and searching a Thing Description Directory exploration service or by fetching a Thing Description directly from the well-known endpoint on a Thing.

**Domain-specific Vocabulary**

Linked Data vocabulary that can be used in the WoT Thing Description, but is not defined by W3C WoT.

**Edge Device**

A device that provides an entry point into enterprise or service provider core networks. Examples include hubs, gateways, routers, switches, multiplexers, and a variety of other access devices.

**Enriched TD**

A Thing Description embedded with additional attributes for bookkeeping and discovery.

**Event**

An Interaction Affordance that describes an event source, which asynchronously pushes event data to Consumers (e.g., overheating alerts).

### ***Exploration***

A discovery mechanism that provides access to detailed metadata in the form of one or more Thing Descriptions. Exploration mechanisms are in general protected by security mechanism and are accessible only to authorized users.

### ***Exposed Thing***

A software abstraction that represents a locally hosted Thing that can be accessed over the network by remote Consumers. The abstraction might be created by a native WoT Runtime, or instantiated as an object through the WoT Scripting API.

### ***Exposing a Thing***

To create an Exposed Thing software abstraction in the local runtime environment to manage the state of a Thing and interface with the behavior implementation.

### ***Hypermedia Control***

A serialization of a Protocol Binding in hypermedia, that is, either a Web link [[RFC8288](#)] for navigation or a Web form for performing other operations. Forms can be seen as request templates provided by the Thing to be completed and sent by the Consumer.

### ***Interaction Affordance***

Metadata of a Thing that shows and describes the possible choices to Consumers, thereby suggesting how Consumers may interact with the Thing. There are many types of potential affordances, but W3C WoT defines three types of Interaction Affordances: Properties, Actions, and Events. A fourth Interaction Affordance is navigation, which is already available on the Web through linking.

### ***Interaction Model***

An intermediate abstraction that formalizes and narrows the mapping from application intent to concrete protocol operations. In W3C WoT, the defined set of Interaction Affordances constitutes the Interaction Model.

### ***Intermediary***

An entity between Consumers and Things that can proxy, augment, or compose Things and republish a WoT Thing Description that points to the WoT Interface on the Intermediary instead of the original Thing. For Consumers, an Intermediary may be indistinguishable from a Thing, following the Layered System constraint of REST.

### ***Introduction***

A "first contact" discovery mechanism, whose result is a URL that references an exploration mechanism. Introduction mechanisms themselves should not directly provide metadata, and in general are designed to be open.

### ***IoT Platform***

A specific IoT ecosystem such as OCF, oneM2M, or Mozilla Project Things with its own specifications for application-facing APIs, data model, and protocols or protocol configurations.

### ***Metadata***

Data that provides a description of an entity's abstract characteristics. For example, a Thing Description is Metadata for a Thing.

### ***Personally Identifiable Information (PII)***

Any information that can be used to identify the natural person to whom such information relates, or is or might be directly or indirectly linked to a natural person. We use the same definition as [\[ISO-IEC-29100\]](#).

### ***Orchestration***

The automation of the behavior of a collection of things. Orchestration is combining individual things with rules or services into a new service or virtual [Thing](#).

### ***Partial TD***

A [Partial TD](#) is an object that follows the same hierarchical structure of the [TD](#) information model, but it is not required to contain all the mandatory elements.

Note: An example for the usage of a [Partial TD](#) is in [WoT Scripting API](#), where it is used as input for the creation of [Exposed Things](#).

### ***Privacy***

Freedom from intrusion into the private life or affairs of an individual when that intrusion results from undue or illegal gathering and use of data about that individual. We use the same definition as [\[ISO-IEC-2382\]](#). See also [Personally Identifiable Information](#) and [Security](#), as well as other related definitions in [\[ISO-IEC-29100\]](#).

### ***Private Security Data***

Private Security Data is that component of a Thing's [Security Configuration](#) that is kept secret and is not shared with other devices or users. An example would be private keys in a PKI system. Ideally such data is stored in a separate memory inaccessible to the application and is only used via abstract operations, such as signing, that do not reveal the secret information even to the application using it.

### ***Producer***

An entity that can create WoT Thing Descriptions for a specific Thing.

### ***Profile***

A technical specification which provides a set of assertions such that any [Consumer](#) which conforms with the those assertions is out-of-the-box interoperable with any [Thing](#) which also conforms with those assertions.

### ***Property***

An Interaction Affordance that exposes state of the Thing. This state can then be retrieved (read) and optionally updated (write). Things can also choose to make Properties observable by notifying [Consumers](#) about a state change.

### ***Protocol Binding***

The mapping from an Interaction Affordance to concrete messages of a specific protocol, thereby informing [Consumers](#) how to activate the [Interaction Affordance](#). ~~W3C~~ WoT serializes Protocol Bindings as [hypermedia controls](#).

### ***Public Security Metadata***

Public Security Metadata is that component of a Thing's [Security Configuration](#) which describes the security mechanisms and access rights necessary to access a Thing. It does not

include any secret information or concrete data (including public keys), and does not by itself, provide access to the Thing. Instead, it describes the mechanisms by which access may be obtained by authorized users, including how they must authenticate themselves.

### ***Registrant***

An entity which registers a Thing Description with a Thing Description Directory. This entity may or may not be the Thing that the registered Thing Description describes.

### ***Security***

Preservation of the confidentiality, integrity and availability of information. Properties such as authenticity, accountability, non-repudiation, and reliability may also be involved. This definition is adapted from the definition of *Information Security* in [ISO-IEC-27000], which also includes additional definitions of each of the more specific properties mentioned. Please refer to this document for other related definitions. We additionally note that it is desirable that these properties be maintained both in normal operation and when the system is subject to attack.

### ***Security Configuration***

The combination of Public Security Metadata, Private Security Data, and any other configuration information (such as public keys) necessary to operationally configure the security mechanisms of a Thing.

### ***Service***

A Service is a software entity that has a network interface. Services can be described by a Thing Description and are a kind of Thing. See also Virtual Thing. Compare with Device.

### ***Servient***

A software stack that implements the WoT building blocks. A Servient can host and expose Things and/or host Consumers that consume Things. Servients can support multiple Protocol Bindings to enable interaction with different IoT platforms.

### ***Shadow***

A Shadow is a Virtual Thing that maintains a copy of the state and mediates interactions with another Thing. A Shadow aims to achieve eventual consistency with the state of the Thing it represents. If a Shadow has more functionality than simply mirroring state it may be better to refer to it as a Digital Twin.

### ***Subprotocol***

An extension mechanism to a transport protocol that must be known to interact successfully. An example is long polling for HTTP.

### ***System***

An entity consisting of multiple interacting components.

### ***TD***

Short for WoT Thing Description.

### ***TDD***

Short for WoT Thing Description Directory.

### ***TD Vocabulary***

A Linked Data vocabulary controlled by [W3C](#) WoT to tag the metadata of Things in the WoT Thing Description including communication metadata of WoT Binding Templates.

### ***TD Context Extension***

A mechanism to extend [Thing Descriptions](#) with additional [Vocabulary Terms](#) using `@context` as specified in JSON-LD[[JSON-LD11](#)]. It is the basis for semantic annotations and extensions to core mechanisms such as Protocol Bindings, Security Schemes, and Data Schemas.

### ***TD Server***

Short for [Thing Description Server](#).

### ***Thing or Web Thing***

An abstraction of a physical or a virtual entity whose metadata and interfaces are described by a WoT Thing Description, whereas a virtual entity is the composition of one or more Things.

### ***Thing Description Directory***

A directory service for TDs that provides a Web interface to register TDs and look them up (e.g., using JSONPath or SPARQL queries). A recommended API and feature set is defined in [[WOT-DISCOVERY](#)], and is used as an optional part of the [WoT Discovery](#) process.

### ***TD Fragment***

A [TD Fragment](#) is a substructure of the data model of a TD. It is a valid object structure that can be validated syntactically against a part of the TD information model defined in chapter 5 of the [Thing Descriptions](#) specification, however the fragment may omit some context that allows full validation.

### ***Thing Description Server***

A Thing Description Server is a web resource, addressed by a URL, that can provide a Thing Description when accessed. Its requirements are defined in [[WOT-DISCOVERY](#)], and is used as an optional part of the [WoT Discovery](#) process.

### ***Thing Model***

A [Thing Model](#) is a description for a class of Things that have the same capabilities. It describes the [Properties](#), [Actions](#), and [Events](#) and common metadata that are shared for an entire group of [Things](#). Compared to a Thing Description, a Thing Model does not contain enough information to identify or interact with a Thing instance.

### ***Transport Protocol***

The underlying, standardized application layer protocol without application-specific requirements or constraints on options or subprotocol mechanisms. Examples are HTTP, CoAP, or MQTT.

### ***Trusted Environment***

Set of devices that assume each other's claims of identity are authentic without proof and allow relatively unrestricted access to one another over a common protected network.

### ***Virtual Thing***

A [Service](#) that represents, augments the functionality of, provides an improved interface to, or stands in place of one or more other [Things](#). A Virtual Thing will often act as an [Intermediary](#). Examples include [Shadows](#) and [Digital Twins](#).

## ***Vocabulary***

A collection of [Vocabulary Terms](#), identified by a namespace IRI.

### ***Term and Vocabulary Term***

A character string. When a [Term](#) is part of a [Vocabulary](#), i.e., prefixed by a namespace IRI[[RFC3987](#)], it is called a [Vocabulary Term](#). For the sake of readability, [Vocabulary Terms](#) present in this document are always written in a compact form and not as full IRIs.

### ***WoT Interface***

The network-facing interface of a Thing that is described by a WoT Thing Description.

### ***WoT Profile***

Synonym for [Profile](#)

### ***WoT Runtime***

A runtime system that maintains an execution environment for applications, and is able to expose and/or consume Things, to process WoT Thing Descriptions, to maintain Security Configurations, and to interface with Protocol Binding implementations. A WoT Runtime may have a custom API or use the optional WoT Scripting API.

### ***WoT Scripting API***

The application-facing programming interface provided by a Servient in order to ease the implementation of behavior or applications running in a WoT Runtime. It is comparable to the Web browser APIs. The WoT Scripting API is an optional building block for [W3C](#) WoT.

### ***WoT Servient***

Synonym for Servient.

### ***WoT Thing Description or Thing Description***

Structured data describing a Thing. A WoT Thing Description comprises general metadata, domain-specific metadata, Interaction Affordances (which include the supported Protocol Bindings), and links to related Things. The WoT Thing Description format is the central building block of [W3C](#) WoT.

## **3.1 Device Categories**

In a deployment of WoT conforming to the [WoT Abstract Architecture](#) we see a variety of different device types. They range (sorted in the order of footprint and capabilities) from small embedded *node* devices to *gateways* or *hubs* to powerful *edge* devices and *cloud* servers. Interoperability between these devices implies that a core set of features and functionalities is available on *all* of them.

The following device categories describe the footprint and characteristics of typical representatives of these classes. This is used to identify the possible features and use cases for these device classes.

These categories are aligned with the classes defined by the IETF [RFC7228] for constrained devices, however the classes have been extended to larger devices and bounds on typical sizes of RAM and Flash/ROM are provided. Memory and storage size are both easier to quantify and more limiting than performance, so that is the basis of this categorization. This is not a strict categorization. Categories may overlap and not all memory may be available for user applications.

<i>Category</i>	<i>data size (RAM)</i>	<i>code size (Flash, ROM, ...)</i>	<i>Typical Representative</i>	<i>Remarks, typical application scenarios</i>
Class-0, C0	<< 10 KiB	<< 100 KiB	small footprint microcontrollers	Sensor nodes without secure communication
Class-1, C1	~ 10 KiB	~ 100 KiB	microcontrollers	Sensors that typically supports secure communication protocols such as TLS, DTLS
Class-2, C2	~ 64 KiB	~ 256 KiB	connected limited device	Small embedded devices such as M2M communication nodes, smart meters, sensor nodes and other embedded appliances, home appliances, low-end TV set-top boxes, and point of sale terminals are some examples.
Class-3, C3	~ 64-256 KiB	~ 256 KiB - several MBs	ISP gateway	Small home and industrial gateways
Class-4, C4	~ 256 KiB - several MB	~ 1 MB - several MB	gateway/hub	Large home and industrial gateways
Class-5, C5	~ 1 - 8 GB	~ 1 - 16 GB	edge	Small edge servers
Class-6, C6	~ several GB	~ several GB	edge	Large edge servers



Class-7, C7	~ several GB	~ several GB	cloud	Cloud systems with multiple compute nodes
----------------	-----------------	-----------------	-------	--

#### NOTE: Category Borders

Category borders are soft borders and categories are not exclusive, i.e. there are application scenarios that can be implemented by devices from several categories. For secure communications we consider devices that support TLS/HTTP, DTLS/CoAP, and similar secure protocols.

## § 4. Application Domains

*This section is non-normative.*

This section presents the application domains targeted by the ~~W3C~~ WoT and which are used to derive the abstract architecture discussed in [7. WoT Building Blocks](#).

These application domains are motivated by the use cases that are described in [[WOT-USE-CASES-REQUIREMENTS](#)].

The Web of Things architecture does not put any limitations on use cases and application domains. Various application domains have been considered to collect common patterns that have to be satisfied by the abstract architecture.

The following sections are not exhaustive. Rather they serve as illustrations, where connected things can provide additional benefit or enable new scenarios.

### § 4.1 Consumer

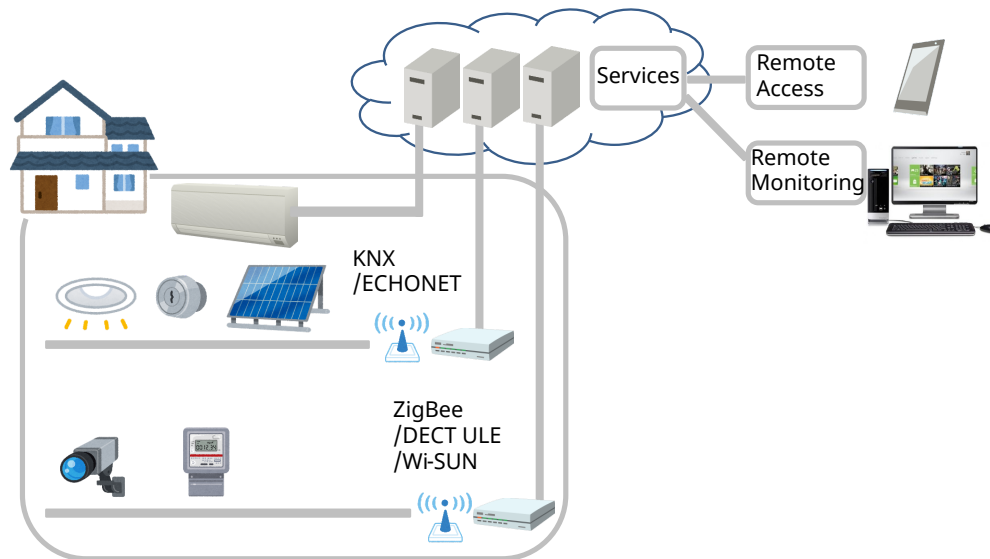
In the consumer space there are multiple assets that benefit from being connected. Lights and air conditioners can be turned off based on room occupancy. Window blinds can be closed automatically based on weather conditions and presence. Energy and other resource consumption can be optimized based on usage patterns and predictions.

The consumer scenario in this section describes the Smart Home use case.

[Figure 1](#) shows an example of a Smart Home. In this case, gateways are connected to edge devices such as sensors, cameras and home appliances through corresponding local communication

protocols such as KNX, ECHONET, ZigBee, DECT ULE and Wi-SUN. Multiple gateways can exist in one home, while each gateway can support multiple local protocols.

Gateways can be connected to the cloud through the internet, while some appliances can be connected to the cloud directly. Services running in the cloud collect data from edge devices and analyze the data, then provide value to users through the edge devices and other UX devices.



*Figure 1 Smart Home*

Smart home provides consumer benefits such as remote access and control, voice control and home automation. Smart home also enables device manufacturers to monitor and maintain devices remotely. Smart home can realize value added services such as energy management and security surveillance.


## § 4.2 Industrial

The industrial use cases in this section are applicable to different industry verticals. Due to the nature of overlaps in the application scenarios, different verticals have similar use cases.

### § 4.2.1 Example: Smart Factory

[Figure 2](#) shows an example of a Smart Factory. In this case, field-level, cell and line controllers automate different factory equipment based on industrial communication protocols such as PROFINET, Modbus, OPC UA TSN, EtherCAT, or CAN. An industrial edge device collects

selected data from various controllers and makes it available to a cloud backend service, e.g., for remote monitoring via a dashboard or analyzes it for preventive maintenance.



*Figure 2 Smart Factory*

Smart factories require advanced monitoring of the connected manufacturing equipment as well of the manufactured products. They benefit from predictions of machine failures and early discovery of anomalies to prevent costly downtime and maintenance efforts.

Additionally, monitoring of connected manufacturing equipment and the environment at the production facility for the presence of poisonous gases, excessive noise or heat increases the safety of the workers and reduces the risks of incidents or accidents.

Real-time monitoring and KPI calculations of production equipment helps to detect productivity problems and optimize the supply chain.

## § 4.3 Transportation & Logistics

Monitoring of vehicles, fuel costs, maintenance needs and assignments helps to optimize the full utilization of the vehicle fleet.

Shipments can be tracked to be en-route to ensure consistent quality and condition of the transported goods. This is especially useful to assert the integrity of the cold-chain from warehouses to refrigerated trucks to delivery.

Centralized monitoring and management of stock in warehouses and yards can prevent out of stock and excessive stock situations.

## § 4.4 Utilities

Automated reading of residential and C&I (Commercial and Industrial) meters, and billing offers continuous insights into resource consumption and potential bottlenecks.

Monitoring the condition and output of distributed renewable energy generation equipment enables optimization of distributed energy resources.

Monitoring and remote-controlling of distribution equipment helps to automate the distribution process.

Continuous monitoring of generation and distribution infrastructure is improving safety of utilities crew in the field.

## § 4.5 Oil and Gas

Offshore platform monitoring, leakage detection and prediction of pipelines as well as monitoring and controlling the levels in tanks and reservoirs helps to improve the industrial safety for the workforce as well as for the environment.

Automated calculation of a distributed stock through various storage tanks and delivery pipes/trucks allows for improved planning and resource optimization.

## § 4.6 Insurance

Proactive Asset Monitoring of high value assets such as connected structures, fleet vehicles, etc. mitigates the risk of severe damage and high costs due to predictions and early detection of incidents.

Usage based insurance can be offered with usage tracking and customized insurance policies.

Predictive weather monitoring and re-routing fleet vehicles to covered garages can limit loss due to hail damage, tree damage.

## § 4.7 Engineering and Construction

Monitoring for industrial safety reduces the risks of security hazards. Monitoring of assets at construction site can prevent damage and loss.

## § 4.8 Agriculture

Soil condition monitoring and creating optimal plans for watering, fertilizing as well as monitoring the produce conditions optimize the quality and output of agricultural produce.

## § 4.9 Healthcare

Data collection and analytics of clinical trial data helps to gain insights into new areas.

Remote patient monitoring mitigates the risk of undetected critical situations for elderly people and patients after hospitalization.

## § 4.10 Environment Monitoring

Environment monitoring typically relies on a lot of distributed sensors that send their measurement data to common gateways, edge devices and cloud services.

Monitoring of air pollution, water pollution and other environmental risk factors such as fine dust, ozone, volatile organic compound, radioactivity, temperature, humidity to detect critical environment conditions can prevent unrecoverable health or environment damages.

## § 4.11 Smart Cities

Monitoring of Bridges, Dams, Levees, Canals for material condition, deterioration, vibrations discovers maintenance repair work and prevents significant damage. Monitoring of highways and providing appropriate signage ensures optimized traffic flow.

Smart Parking is optimizing and tracking the usage and availability of parking spaces and automates billing/reservations.

Smart control of street lights based on presence detection, weather predictions, etc. reduces cost.

Garbage containers can be monitored to optimize the waste management and the trash collection route.

## § 4.12 Smart Buildings

Monitoring the energy usage throughout the building helps to optimize resource consumption and reduce waste.

Monitoring the equipment in the buildings such as HVAC, Elevators, etc. and fixing problems early improves the satisfaction of occupants.

## § 4.13 Connected Car

Monitoring of operation status, prediction of service needs optimizes maintenance needs and costs. Driver safety is enhanced with notifications of an early warning system for critical road and traffic conditions.

### § 4.13.1 Connected Car Example

[Figure 3](#) shows an example of a Connected Car. In this case, a gateway connects to car components through CAN and to the car navigation system through a proprietary interface. Services running in the cloud collect data pushed from car components and analyze the data from multiple cars to determine traffic patterns. The gateway can also consume cloud services, in this case, to get traffic data and show it to the driver through the car navigation system.

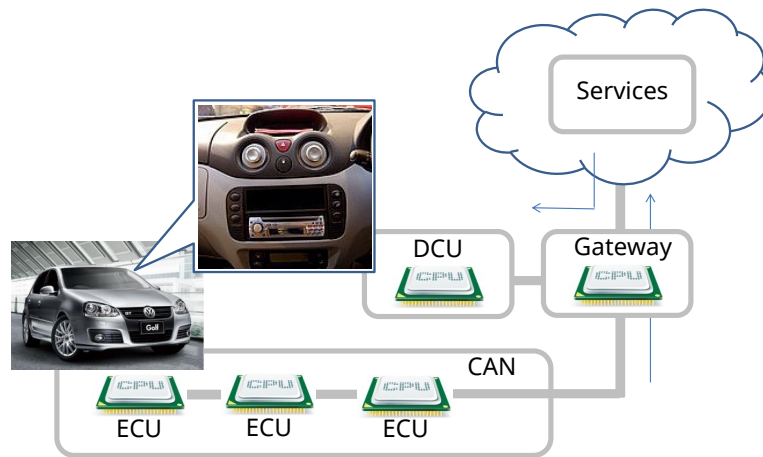


Figure 3 Connected Car

## § 5. Common Deployment Patterns

*This section is non-normative.*

This section introduces common deployment patterns that illustrate how devices/things interact with controllers, other devices, agents and servers. In this section, we use the term *client role* as an initiator of a transport protocol, and the term *server role* as a passive component of a transport protocol. This does not imply prescribing a specific role on any system component. A device can be in a *client* and *server* role simultaneously.

One example of this dual role is a sensor, that registers itself with a cloud service and regularly sends sensor readings to the cloud. In the response messages the cloud can adjust the transmission rate of the sensor's messages or select specific sensor attributes, that are to be transmitted in future messages. Since the sensor registers itself with the cloud and initiates connections, it is in the 'client' role. However, since it also reacts to requests, that are transmitted in response messages, it also fulfills a 'server' role.

The following sections illustrate the roles, tasks, and use case patterns with increasing complexity. They are not exhaustive and are presented to motivate for the WoT architecture and building blocks that are defined in later sections of this specification.

This section also makes use of the concept of a Trusted Environment, which is a set of devices that allow relatively unrestricted access to one another. This is a common approach but carries some risks, which are discussed in section [10.4 Trusted Environment Risks](#), along with mitigations of these risks.

## § 5.1 Telemetry

Telemetry is the monitoring of remote devices, which implies automatic transmission to, and processing of measurements and other data on a Consumer. Data may be transported over various communication channels, both wireless and wired. Examples include GSM networks, Bluetooth, WiFi, Ethernet, and other wired standards.

Remote metering devices include one or multiple sensors, in some cases they also contain an actuator. In many cases the sensor data is transmitted at regular intervals, on a state change, or as a response to a request by the Consumer.

Remote metering devices are frequently small embedded systems with very limited resources, i.e. Class-2 or below. They may be battery powered and have to minimize power consumption through various measures, e.g. sleep modes. These devices will sleep most of the time and won't transmit any data to conserve energy. They only wake up on certain events (e.g. when a switch is toggled). When a device state changes, an event is sent to the consumer. After that the device goes to sleep mode again.

Typical user scenarios are monitoring of various assets, examples include smart cities, factories, fleets, environmental monitoring and health monitoring. Some deployments may require monitoring of a high number of geographically distributed assets, others only include a few devices, as in a smart home. A common pattern is the one-to-many relationship between a single Consumer and multiple devices (Things).


## § 5.2 Device Controllers

A common deployment pattern is a local device controlled by a user-operated remote controller as depicted in [Figure 4](#).

A remote controller can access an electronic appliance through the local home network directly. In this case, the remote controller can be implemented by a browser or native application.

In this pattern, at least one device like the electronic appliance has a server role that can accept a request from the other devices and responds to them, and sometimes initiates a mechanical action. The other device like the remote controller has a client role that can send a message with a request, like to read a sensor value or to turn on the device. Moreover, to emit a current state or event notification of a device, the device may have a client role that can send a message to another device, which has server roles.






*Figure 4 Device Control*

### § 5.3 Thing-to-Thing

[Figure 5](#) shows an example of a direct Thing-to-Thing interaction. The scenario is as follows: a sensor detects a change of the room condition, for example the temperature exceeding a threshold, and issues a control message like "turn on" to the electronic appliance. The sensor unit can issue some trigger messages to other devices.

In this case, when two devices that have server roles are connected, at least one device must have also a client role that issues a message to the other to actuate or notify.



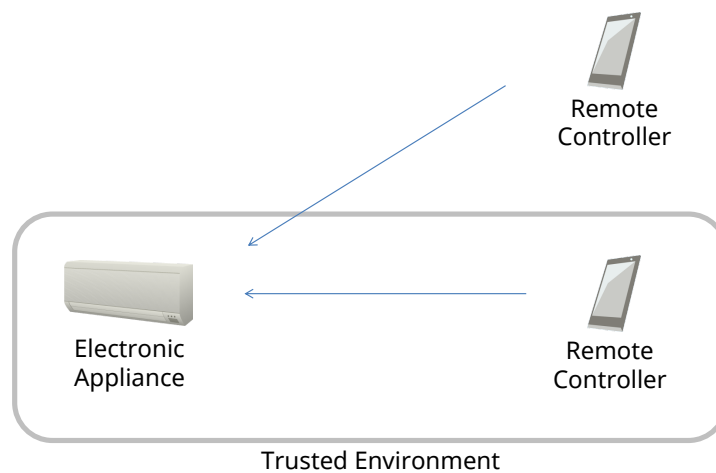
*Figure 5 Control Agent*

### § 5.4 Remote Access

This deployment scenario contains a mobile remote controller (e.g., on a smartphone) as shown in [Figure 6](#). The remote controller can switch between different network connections and protocols, e.g., between a cellular network and a home network, which is using protocols such as Wi-Fi and Bluetooth. When the controller is in the home network it is a trusted device and no additional security or access control is required. When it is outside of the trusted network, additional access control and security mechanisms must be applied to ensure a trusted relationship. Note that in this

scenario the network connectivity may change due to switching between different network access points or cellular base stations.

In this pattern, the remote controller and the electronic appliance have a client and a server role as in the related scenario in [Figure 4](#).




*Figure 6 Multiple Network Interfaces*

## § 5.5 Smart Home Gateways

[Figure 7](#) shows a deployment scenario of a *Smart Home Gateway*. The gateway is placed between a home network and the Internet. It manages electronic appliances inside the house and can receive commands from a remote controller over the Internet, e.g., from a smartphone as in the previous scenario. It is also a virtual representation of a device. The Smart Home Gateway typically offers proxy and firewall functionality.


In this pattern, the home gateway has both a client and a server role. When the remote controller actuates the electronic appliance, it can connect to the electronic appliance in the client role and to the remote controller with the server role. When the electronic appliance emits a message to the remote controller, the gateway acts as server roles for the electric appliance, and it acts as client roles for the remote controller.



*Figure 7 Smart Home Gateway*

## § 5.6 Edge Devices


An *Edge Device* or *Edge Gateway* is similar to a Smart Home Gateway. We use the term to indicate additional tasks that are carried out by the edge gateway. Whereas the home gateway in [Figure 8](#) primarily just bridges between the public and the trusted network, the edge device has local compute capabilities and typically bridges between different protocols. Edge devices are typically used in industrial solutions, where they can provide preprocessing, filtering and aggregation of data provided by connected devices and sensors.



*Figure 8 Edge device*


## § 5.7 Digital Twins

A *Digital Twin* is a virtual representation, i.e. a model of a device or a group of devices that resides on a cloud server or edge device. It can be used to represent real-world devices which may not be continuously online, or to run simulations of new applications and services, before they get deployed to the real devices.



*Figure 9 Digital Twin*

Digital twins can model a single device, or they can aggregate multiple devices in a virtual representation of the combined devices.




*Figure 10 Digital Twin for Multiple Devices*

Digital twins can be realized in different ways, depending on whether a device is already connected to the cloud, or whether it is connected to a gateway, which itself is connected to the cloud.

### § 5.7.1 Cloud Connected Devices

[Figure 11](#) shows an example where electronic appliances are connected directly to the cloud. The cloud mirrors the appliances and, acting as a digital twin, can receive commands from remote controllers (e.g., a smartphone). Authorized controllers can be located anywhere, as the digital twin is globally reachable.



*Figure 11 Appliance Twin for a Cloud Connected Devices*

### § 5.7.2 Orchestration

Orchestration of devices and operation flows provides capabilities that go beyond the simple isolated usage of a device. An orchestrated set of devices can provide combined operations, that affect the state of multiple devices with a single operation.

These operations combine individual operations into a combined operation flow, the flow may have several alternative paths, that are chosen based on some device state or property values.

Orchestrated devices can be deployed in various topologies and can be connected via different networks. The system, that is built from these devices, can provide capabilities, that combine and exceed the capabilities of individual devices.


A typical orchestration example is a smart home, where various sensors (temperature, occupancy, humidity, light, air quality, door sensors) are interworking and provide operation flows for situations like entering or leaving the house, morning/evening routine, adjusting light and temperature based on presence, and much more.

### § 5.7.3 Legacy Devices

[Figure 12](#) shows an example where legacy electronic appliances are not directly connected to the cloud. Here, a gateway is needed to relay the connection. The gateway works as:

- integrator of a variety of legacy communication protocols both in the physical and logical view
- firewall toward the Internet
- privacy filter which substitutes real image and/or speech, and logs data locally
- local agent in case the network connection is interrupted
- emergency services running locally when fire alarms and similar events occur

The cloud mirrors the gateway with all connected appliances and acts as a digital twin that manages them in the cloud in conjunction with the gateway. Furthermore, the cloud can receive commands from remote controllers (e.g., a smartphone), which can be located anywhere.



*Figure 12 A Digital Twin for a Legacy Device*

### § 5.8 Multi-Cloud

Typical IoT deployments consist of multiple (thousands) of devices. Without a standardized mechanism, the management of firmware updates for specific clouds require a lot of effort and hinders wider scale IoT adoption.

The primary benefit of a standardized mechanism for describing devices and device types is the capability of deploying devices to different cloud environments without the need of doing customization at device software / firmware level, i.e., installing cloud specific code to a device. This implies that the solution is flexible enough to describe devices in a way that allows on-boarding and using devices in multiple IoT cloud environments.

This drives adoption of Web of Things devices, since it enables easy usage of new devices in an existing deployment, as well as migration of existing devices from one cloud to the other.

## § 5.9 Virtual Things

A Virtual Thing is a Service that acts as a placeholder for one or more other Things which provides an interface to its consumers.


In a simple case it is an interface abstraction, which defines a common interface to devices, where it is possible to replace a device model with another one without changing the consumer.

In a more complex scenario, a Virtual Thing provides a single interface for multiple devices, and provides a higher level of operations to its Consumers. Individual devices can be replaced, new operations can be provided via software upgrades.

A Virtual Thing will often act as an Intermediary. Examples include Shadows and Digital Twins.

## § 5.10 Cross-domain Collaboration

Figure 13 show an example of a cross-domain collaboration. In this case, each system involves other systems in other domains, such as Smart Factory with Smart City, Smart City with Smart Home. This type of system is called "Symbiotic" ecosystem, as shown in [IEC-FOTF]. There are two collaboration models: direct collaboration and indirect collaboration. In the direct collaboration model, systems exchange information directly with each other in a peer-to-peer manner. In the indirect collaboration, systems exchange information via some collaboration platform. In order to maintain and continue this collaboration, each system provides the metadata of their capabilities and interfaces and adapts itself to others.



*Figure 13 Cross-domain collaboration*

## § 5.11 System Integration

The previous section described various architecture patterns. In these patterns, some functional entities such as the devices including the legacy devices, controllers, gateways and cloud servers are located at physical locations such as inside building, outside buildings, and data centers. [Figure 14](#) is an overview that shows the combinations and communication paths of these entities.

In a transport protocol layer, each entity arbitrarily selects a suitable role for communications. For example, a device may act as a server when the device provides a service to indefinite number of applications. On the other hand, if a device has limited or intermittent network connectivity, they may act as a client and actively send message to an application when network is available. Regardless of this, in application layer, an application sees that a device provides abstract interfaces to interact and the application can interact with the device using their abstract interfaces.






Figure 14 System integration

## § 6. Abstract WoT System Architecture

*This section is normative.*

To address the requirements and use cases that were gathered in [[WOT-USE-CASES-REQUIREMENTS](#)] the Web of Things (WoT) builds on top of the concept of Web Things – usually simply called Things – that can be used by so-called Consumers. Consumers can interact with a Thing directly, or they use intermediaries for indirect communication.

These concepts are applicable for the various application domains of section 4. [Application Domains](#).

This section provides the background and normative assertions to define the overall ~~W3C~~ Web of Things architecture.

As the Web of Things addresses stakeholders from different domains, certain aspects of Web technology are explained in more detail, in particular the concept of hypermedia.

### § 6.1 Fundamental Concepts

This section introduces the fundamental concepts of the Web of Things. It introduces the Thing abstraction that is described with Thing Descriptions and Thing Models, which then can be used by

## Consumers.

A Thing is an abstraction of a physical or virtual entity (e.g., a device or a room) and is described by standardized metadata. A Consumer is an entity that can read and interpret the standardized metadata for a Thing in order to communicate with that Thing.

The WoT Thing Description (TD) is a standardized, machine-readable metadata representation format that allows Consumers to discover and interpret the capabilities of a Thing (through semantic annotations) and to adapt to different implementations (e.g., different protocols or data structures) when interacting with a Thing, thereby enabling interoperability across different IoT platforms, i.e., different ecosystems and standards. The WoT Thing Model (TM) is likewise a standardized, machine-readable metadata representation format for describing *classes* of Things, such as devices of a product line with common capabilities.



Figure 15 Consumer-Thing interaction

A Thing can also be the abstraction of a *virtual entity*. A virtual entity is the composition of one or more Things (e.g., a room consisting of several sensors and actuators). One option for the composition is to provide a single, consolidated WoT Thing Description that contains the combination of capabilities for the virtual entity. In cases where the composition is rather complex, its TD may *link* to hierarchical sub-Things within the composition. The main TD acts as entry point and only contain general metadata and potentially overarching capabilities. This allows grouping of certain aspects of more complex Things.

### § 6.1.1 Metadata

The WoT architecture provides metadata formats to describe both specific *instances* of Things and *classes* of Things. The metadata format for *instances* is called Thing Description while that for *classes* is called Thing Model.

#### § 6.1.1.1 Thing Descriptions

A Thing instance is described by standardized metadata. In W3C WoT, the description metadata for a Thing instance *MUST* be available as a WoT Thing Description (TD) [WOT-THING-DESCRIPTION]. The format can be processed either through classic JSON libraries or a JSON-LD processor, as the underlying information model is graph-based and its serialization compatible with JSON-LD 1.1 [JSON-LD11]. The use of a JSON-LD processor for processing a TD additionally enables semantic processing including transformation to RDF triples, semantic inference and accomplishing tasks given based on ontological terms, which would make Consumers behave more autonomous. A TD is instance-specific (i.e., describes an individual Thing, not types of Things) and is the default external, textual (Web) representation of a Thing. There *MAY* be other representations of a Thing such as an HTML-based user interface, simply an image of the physical entity, or even non-Web representations in closed systems. To be considered a Thing, however, at least one TD representation *MUST* be available.

#### § 6.1.1.2 *Thing Models*

A Thing Model can be used to describe common capabilities that are available for a set of Things, for example, for a large number of devices in a product line. In this case the Thing Model defines the common capabilities of all devices in the product line, but omits information specific to a particular device.

A Thing Model may also be used when complete information for an instance is not available or is not necessary. For example, some IoT ecosystems implicitly handle communication separately. In such a case, a fully detailed Thing Description, e.g. with communication metadata, is not necessary. Communication metadata may also not be available at the beginning of a Thing lifecycle phase, since a new entity has not yet been deployed (e.g. the IP address is not yet known).

A Thing Model is used to define the basic information model of a Thing to address such kind of scenarios. The Thing Model can be seen as a template for Thing Descriptions, however, that have less restriction as defined in sections "TD Information Model" and "Representation Format" of the [WOT-THING-DESCRIPTION]. Typically Thing Model examples do not contain any instance-specific information such as protocol specific data like IP addresses. However, instead of having, e.g., concrete URLs, Thing Model allows the usage of URL templates.

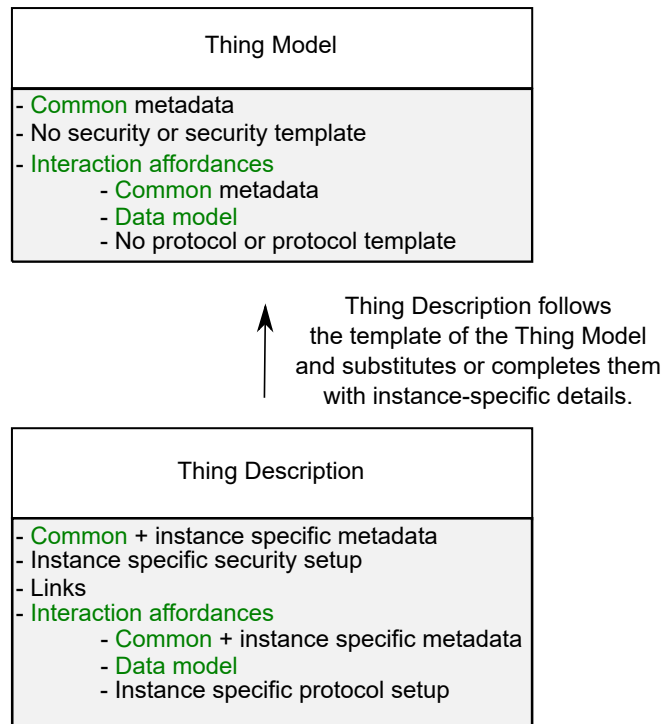


Figure 16 Thing Model and its relation to the Thing Description

The Thing Model enables:

- onboarding and management of multiple Thing models, e.g., by a cloud service.
- simulation of devices/Things that have not yet been developed.
- developing common applications across devices from different manufacturers that share a common Thing model.
- combining multiple models into a Thing.
- implementation support of a concrete Thing.

The Thing Model is a logical description of the interface and possible interaction with Thing's Properties, Actions, and Events, however it does not contain Thing instance-specific information, such as concrete protocol usage (e.g., IP address), or even a serial number and GPS location. However, Thing Models allows to include, e.g., security schemes if they apply to the entire class of instances the model describes. They might have URLs (e.g., like token servers) that might need to be omitted or parameterized (with templates) although in a lot of cases these might also be given.

A Thing Model can be serialized in the same JSON-based format as a Thing Description which also allows JSON-LD processing. Note that a Thing Model cannot be validated in the same way as Thing Description instances, since not all mandatory terms of a Thing Description are required for a Thing Model. missing mandatory terms.

## § 6.1.2 Links

Links can represent relationships between things, between things and thing models, and between thing models. Linking does not only apply to hierarchical Things, but also to relations between Things and other resources in general. Link relation types express how Things relate, for instance, a switch controlling a light or a room monitored by a motion sensor. Other resources related to a Thing can be documentation manuals, catalogs for spare parts, CAD files, a graphical UI, or any other document on the Web. Overall, Web linking among Things makes the Web of Things navigable, for both humans and machines. This can be further facilitated by providing Thing Description Directories that manage a catalog of Things, usually by caching their TD representation.

In summary, WoT Thing Descriptions and WoT Thing Models *MAY* link to other Things, WoT Thing Models, and other resources on the Web to form a Web of Things.




Figure 17 Linked Things


Things *MUST* be hosted on networked system components with a software stack to realize interaction through a network-facing interface, the WoT Interface of a Thing. One example of this is an HTTP server running on an embedded device with sensors and actuators interfacing the physical entity behind the Thing abstraction. However, W3C WoT does not mandate where Things are hosted; it can be on the IoT device directly, an Edge device such as a gateway, or the cloud.

A typical deployment challenge is a scenario where local networks are not reachable from the Internet, usually because of IPv4 Network Address Translation (NAT) or firewall devices. To remedy this situation, W3C WoT allows for Intermediaries between Things and Consumers.

## § 6.1.3 Intermediaries

Intermediaries can act as proxies for Things, where the Intermediary has a WoT Thing Description similar to the original Thing, but which points to the WoT Interface provided by the Intermediary.

Intermediaries may also augment existing Things with additional capabilities or compose a new Thing out of multiple available Things, thereby forming a Virtual Thing. To Consumers, Intermediaries are just another kind of Thing, as they possess WoT Thing Descriptions and provide a WoT Interface. They may be indistinguishable from Things that are directly representing physical devices in a layered system architecture like the Web [REST].



*Figure 18 Intermediary*

Another remedy for restricted local networks is binding the WoT Interface to a protocol that establishes the connection from the Thing within the local network to a publicly reachable Consumer.

The concepts of W3C WoT are applicable to all levels relevant for IoT applications: the device level, edge level, and cloud level. This fosters common interfaces and APIs across the different levels and enables various integration patterns such as Thing-to-Thing, Thing-to-Gateway, Thing-to-Cloud, Gateway-to-Cloud, and even cloud federation, i.e., interconnecting cloud computing environments of two or more service providers, for IoT applications. [Figure 19](#) gives an overview how the WoT concepts introduced above can be applied and combined to address the use cases described in the *WoT Use Cases and Requirements* document [[WOT-USE-CASES-REQUIREMENTS](#)].

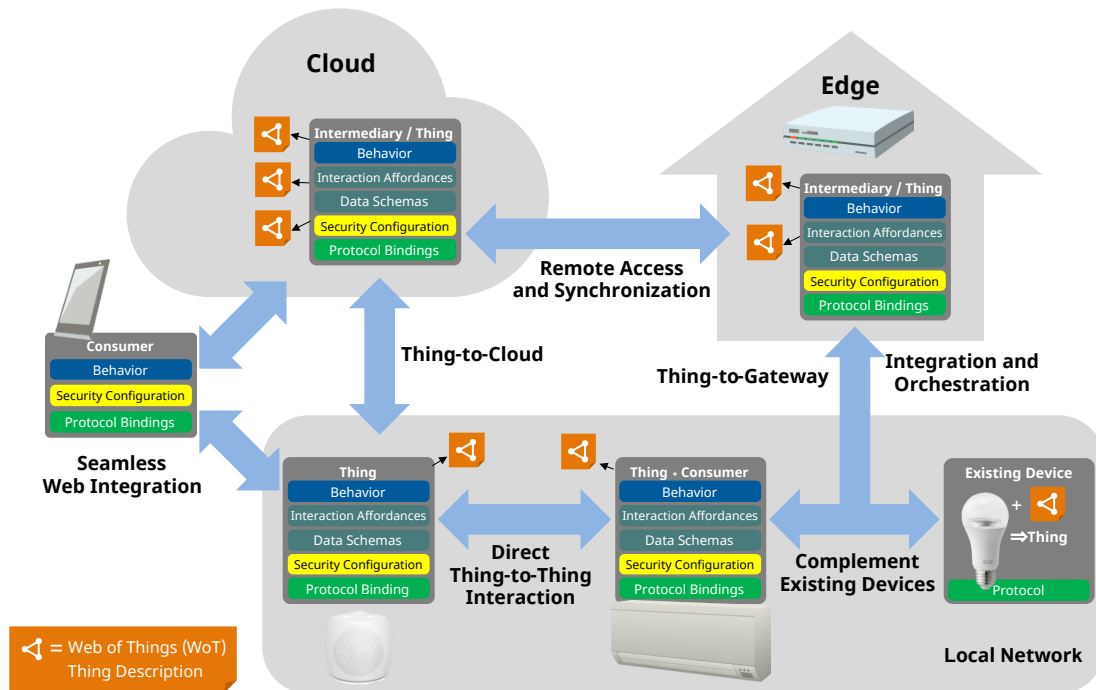


Figure 19 Abstract Architecture of W3C WoT

## § 6.2 Affordances

A central aspect in W3C WoT is the provision of machine-readable metadata (i.e., WoT Thing Descriptions). Ideally, such metadata is self-descriptive, so that Consumers are able to identify *what* capabilities a Thing provides and *how* to use the provided capabilities. A key to this self-descriptiveness lies in the concept of *affordances*.

The term affordance originates in ecological psychology, but was adopted in the field of Human-Computer Interaction [HCI] based on the definition by Donald Norman: "'Affordance' refers to the perceived and actual properties of the thing, primarily those fundamental properties that determine just how the thing could possibly be used." [NORMAN]

An example for this is a door with a handle. The door handle is an affordance, which suggests that the door can be opened. For humans, a door handle usually also suggests *how* the door can be opened; an American knob suggests twisting, a European lever handle suggests pressing down.

The hypermedia principle, which is one of the core foundations of the REST architectural style [REST], demands that any piece of information available on the Web be linked to other pieces of information so that the consumer of the information gets explicit knowledge about how to navigate the Web and control Web applications. Here, the simultaneous presentation of information and control (provided in the form of hyperlinks) is a mechanism that *affords* Web clients the means to drive Web applications. In this context, an affordance is the description of a hyperlink (e.g., via a

link relation type and link target attributes) suggesting Web clients how to navigate and possibly how to act on the linked resource. Hence, links provide navigation affordances.

Drawn from this hypermedia principle, the Web of Things defines Interaction Affordances as metadata of a Thing that shows and describes the possible choices to Consumers, thereby suggesting how Consumers may interact with the Thing. A general Interaction Affordance is navigation, which is activated by following a link, thereby enabling Consumers to browse the Web of Things. 6.5 Interaction Model defines three more types of Interaction Affordances for WoT: Properties, Actions, and Events.

Overall, this WoT definition is aligned with HCI and interaction designers, who create physical Things, as well as the REST and microservice community, who is working on Web services in general.

## § 6.3 Web Thing

A Web Thing has four architectural aspects of interest: its *behavior*, its Interaction Affordances, its security configuration, and its Protocol Bindings, as depicted in [Figure 20](#). The behavior aspect of a Thing includes both the autonomous behavior and the handlers for the Interaction Affordances. The Interaction Affordances provide a model of how Consumers can interact with the Thing through abstract operations, but without reference to a specific network protocol or data encoding. The protocol binding adds the additional detail needed to map each Interaction Affordance to concrete messages of a certain protocol. In general, different concrete protocols may be used to support different subsets of Interaction Affordances, even within a single Thing. The security configuration aspect of a Thing represents the mechanisms used to control access to the Interaction Affordances and the management of related Public Security Metadata and Private Security Data.

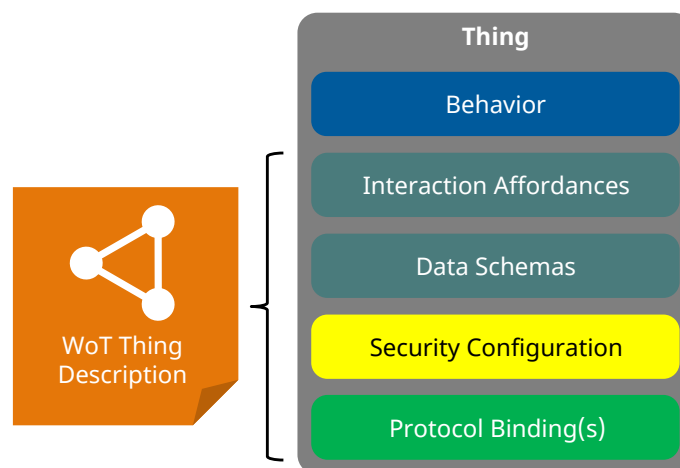


Figure 20 Architectural Aspects of a Thing



## § 6.4 Lifecycle

In a deployment, that applies the WoT architecture principles, different entities interact with each other and exchange information between them. Elements of a WoT deployment are devices, intermediaries, consumers and *directories*. Each element has its own (intrinsic) Thing lifecycle. These elements build a system, where the entire system has a *System lifecycle*, i.e. the system has several states and transitions through some states to become operational. This implies that devices go through their lifecycle to become operational, and devices must be known to others before they can be used.

The following sections describe the *System lifecycle* and the *Thing lifecycle*. They contain sample flows to illustrate the different states and phases of Things in a deployment. One of the purposes of this section is to clarify terminology and ensure that the concepts of a common lifecycle model are considered by implementers of Web Things.

The actor in the left hand side of these diagrams can be understood as the device owner or manufacturer. The actor on the right is the end user, who uses the device in an application.

### § 6.4.1 Simple System Lifecycle

The following sequence diagram shows an example of the flow of system that consists of two directly communicating devices.

Before a device can be used in this scenario, it has to be bootstrapped and onboarded to the consumer. In the onboarding operation the consumer and the device get to know each other - this may either be performed by a discovery process, or via a registration operation, where the consumer is registered with the device or vice-versa.

After an activation step (which may consist of several operations, such as provisioning, configuration etc.) the device and consumer are in regular operation mode. When the device is no longer required in this scenario, it will be offboarded from the consumer. After this operation it can be decommissioned and destroyed.

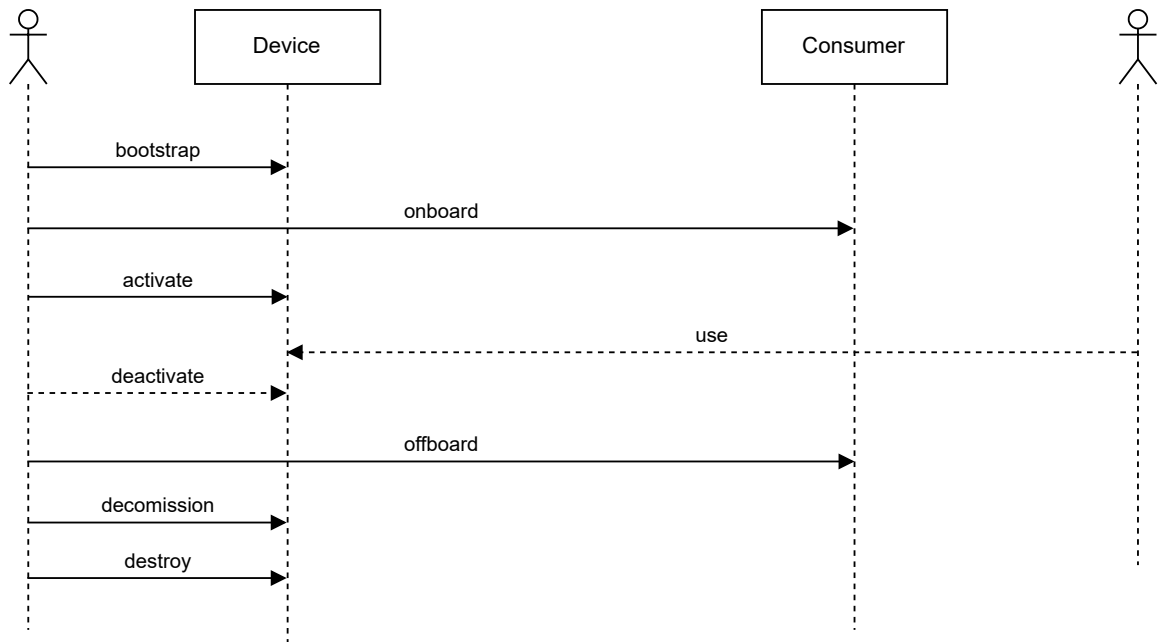


Figure 21 Simple System Lifecycle

#### § 6.4.2 System Lifecycle with Registration

The following sequence diagram shows an example of the flow of system that contains three entities: a device, a consumer and a directory.

This flow is very similar to the flow in the previous section, in addition it contains a *directory* entity, which maintains a catalog of devices. In WoT this catalog is a set of thing descriptions. After a device is bootstrapped like in the previous scenario, it is registered with the directory.

The directory acts as an information broker between the actor on the left and the actor on the consumer side. This decouples the device owner from the consumer, where the discovery acts as a middle man to enable discovery for consumers.

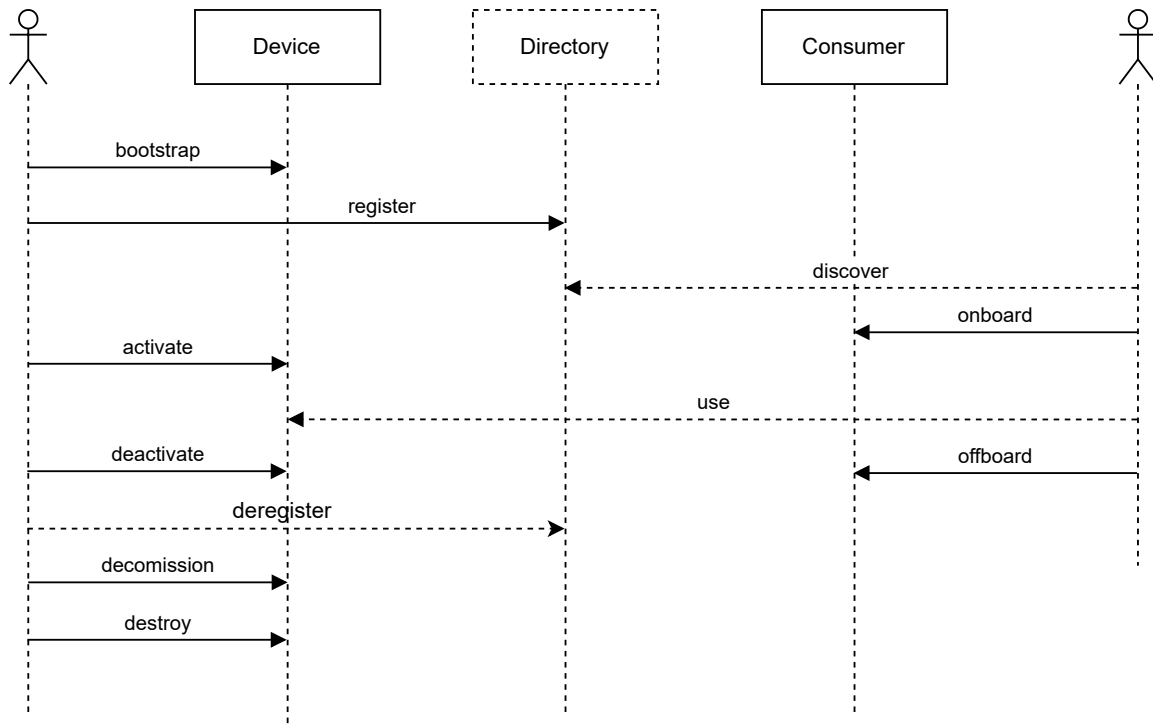


Figure 22 System lifecycle with registration

### § 6.4.3 Thing Lifecycle

Bootstrapping and provisioning of devices is an essential part of setting up devices in all IoT protocol suites.

The main scenarios for provisioning devices with WoT are as follows:

- A device is already provisioned and operational in a given deployment. Make it work with WoT.
- A device is already provisioned and operational in a given deployment. For management purposes, *describe* the device lifecycle stages in a [Thing Description](#).
- Bootstrap and provision a device directly with WoT, in order to become operational for WoT.

Various provisioning schemes are being used in IoT protocol suites. The text in this section is based on [proposals](#) and studies, [comparing various provisioning schemes](#), such as OCF, OneM2M, Lightweight OneM2M, Thing to Thing Research Group (T2TRG), OPC-UA/Anima, etc.

#### NOTE

The provisioning model presented in this section resembles the T2TRG provisioning model.

Common elements of device bootstrapping and provisioning across various IoT protocol suites are as follows:

- Establish the chain of trust, e.g. secure storage, keys, certificates. This may involve various solutions, such as manufacturer certificates, out-of-band key provisioning, connecting to a provisioning server, etc.
- Establish device ownership, using a provisioning tool or service. For instance, the device can be owned by a network entity, or network service, or service provider, or the end user.
- Provision the device with the access control lists for the tenants or various levels of users.
- Provision the device with access to the services it uses.
- Configure the device with used and the exposed services.
- Provision and configure the WoT runtime in the device.
- Update the configurations or provisioning data.
- Decommission a user, application, service, or provisioning.
- Return the device to the initial state before provisioning (e.g. factory reset).
- Decommission and irreversibly destroy a device.

Taken into account these provisioning flows, in general a device can be in one of the following states:

- **Manufactured:** the device is flashed with a software image. In the case it is certified for a certain protocol suite, it may be permitted or capable of doing only limited operations, such as a certain bootstrapping procedure.
- **Bootstrapped:** the device has an identity and ownership established, being ready for the next provisioning steps, like configuration, service provisioning etc. This state has different names in various protocol suites, for instance it is called *onboarded* in OCF, *bootstrapped* in T2TRG, OPC-UA, Anima, LwM2M, *initial provisioning* in OneM2M, etc.
- **Operational:** the device is provisioned and configured, working in normal mode. Some configuration is possible without leaving this state. That may include installing and uninstalling applications, reconfiguring settings, etc. Note that a device can be operational in its own native protocol suite and managed by a WoT gateway, or can be operational for WoT (which may be an application on the device), or may be operational for WoT and directly provisioned for WoT.
- **Maintenance:** the device operational state is interrupted for updating its software and/or configuration.
- **Destroyed:** the device has been wiped out of all data and software. Hardware kill features may be activated. The device may be physically destroyed and never used again. This state is

relevant for device management purposes. It does not exist in OneM2M, LwM2M, OCF, T2TRG and is called *End-of-life* in OPC-UA and Anima.

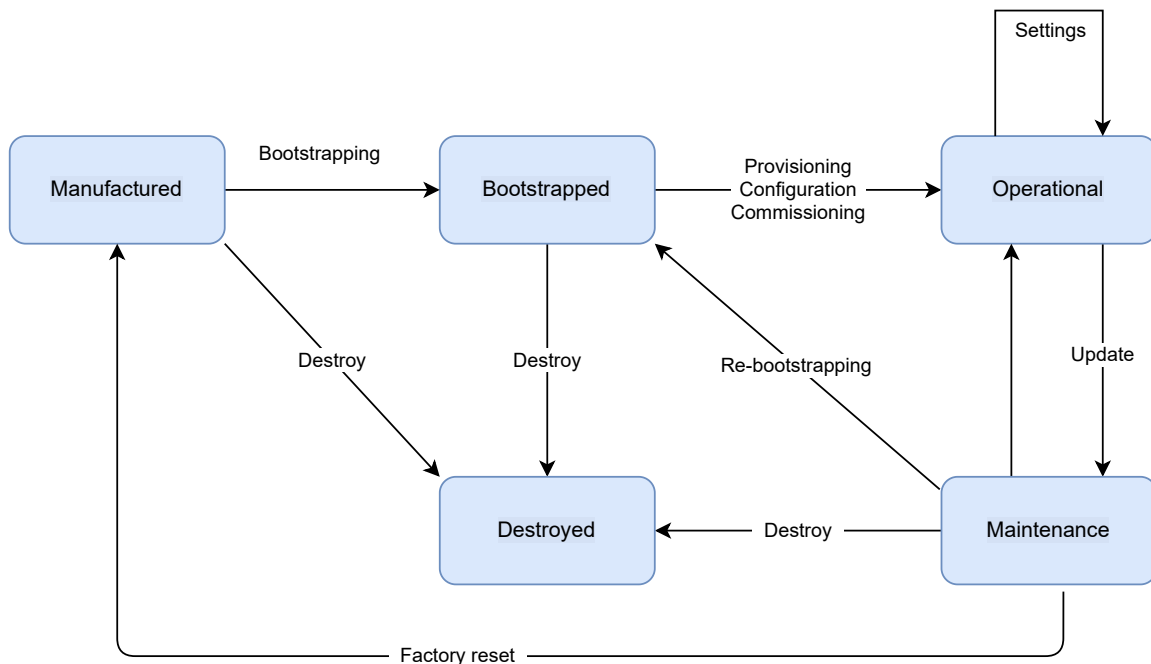


Figure 23 Device Lifecycle

The typical transitions between lifecycle states are the following:

- **Bootstrapping** (or onboarding): the device is provisioned with an identity, the chain of trust is established, e.g. secure storage, keys, certificates. This may involve various solutions, such as manufacturer certificates, out-of-band key provisioning, connecting to a provisioning server (a quite common scenario). When provisioning directly to WoT, the device may be registered with a Thing Description Directory in this stage. During or after this process, in some protocol suites rebooting in a different mode of operation might be also needed.
- **Provisioning, configuration, commissioning**: the device is provisioned with all resources needed for its operation (services, applications, access control, databases etc), and these resources are configured for operation. Also, the device may be commissioned in a given environment. These may involve communication with a server, for instance a Thing Description Directory or Discovery services.
- **Settings**: the device remains in *Operational* state, but may update system, service or application settings. In some cases, this may include installing and removing applications.
- **Update**: the device stops normal operation for undergoing updates in the *Maintenance* state, similar to the ones during provisioning. This may include installing new software, or removing, installing or updating resources and their configuration. It may also include re-commissioning. Returning to *Operational* state may be achieved by resuming operation with updated resources, or may require a restart of services, or rebooting the device.

- **Re-bootstrapping:** the device identity, owner and related resources may be changed as described in the *Bootstrapping* process.
- **Factory Reset:** the device is returned to its factory default state.
- **Destroy:** the device is erased from all data, software and may be physically destroyed.

## § 6.5 Interaction Model

Originally, a Web resource usually represented a document on the World Wide Web that can simply be fetched by a Web client. With the introduction of Web services, resources became more generic interaction entities that can implement any kind of behavior. This very high level of abstraction makes it hard to provide a loose coupling between applications and resources due to the manifold interaction possibilities. As a result, at the time of writing, typical API descriptions consist of a static mapping from an application intent to a resource address, method, request payload structure, response payload structure, and expected errors. This imposes a tight coupling between Web client and Web service.

The Interaction Model of W3C WoT introduces an intermediate abstraction that formalizes the mapping from application intent to concrete protocol operations and also narrows the possibilities how Interaction Affordances can be modeled.

In addition to navigation affordances (i.e., Web links), Things *MAY* offer three other types of Interaction Affordances defined by this specification: Properties, Actions, and Events. While this narrow waist allows to decouple Consumers and Things, these four types of Interaction Affordances are still able to model virtually all interaction kinds found in IoT devices and services.

### § 6.5.1 Properties

A Property is an Interaction Affordance that exposes the state of the Thing. The state exposed by a Property can be retrievable (readable). Optionally, the state exposed by a Property may be updated (writable). Things may choose to make Properties observable by pushing the new state after a change (cf. Observing Resources [RFC7641]). In the case of a read-only state, an Action can be used to update the state.

If the data format is not fully specified by the Protocol Binding used (e.g., through a media type), Properties *MAY* contain one data schema for the exposed state.

Examples of Properties are sensor values (read-only), stateful actuators (read-write), configuration parameters (read-write), Thing status (read-only or read-write), or computation results (read-only).

## § 6.5.2 Actions

An Action is an Interaction Affordance that allows to invoke a function of the Thing. An Action *MAY* manipulate state that is not directly exposed (cf. Properties), manipulate multiple Properties at a time, or manipulate Properties based on internal logic (e.g., toggle). Invoking an Action *MAY* also trigger a process on the Thing that manipulates state (including physical state through actuators) over time.

If the data format is not fully specified by the Protocol Binding used (e.g., through a media type), Actions *MAY* contain data schemas for input parameters and output results.

Some examples for Actions are setting multiple Properties simultaneously, changing Properties over time such as fading the brightness of a light (dimming) or with a process that shall not be disclosed such as a proprietary control loop algorithm, or invoking a long-lasting process such as printing a document.

## § 6.5.3 Events

An Event is an Interaction Affordance, which describes an event source that pushes data asynchronously from the Thing to the Consumer. Here not state, but state transitions (i.e., events) are communicated. Events *MAY* be triggered through conditions that are not exposed as Properties.

If the data is not fully specified by the Protocol Binding used (e.g., through a media type), Events *MAY* contain data schemas for the event data and subscription control messages (e.g., a callback URI to subscribe with a Webhook).

Examples of Events are discrete events such as an alarm or samples of a time series that are sent regularly.

## § 6.6 Hypermedia Controls

On the Web, an affordance is the simultaneous presentation of information and controls, such that the information becomes the affordance through which the user obtains choices. For humans, the information is usually text or images describing or decorating a hyperlink. The control is a Web link, which includes at least the URI of the target resource, which can be dereferenced by the Web browser (i.e., the link can be followed). But also machines can follow links in a meaningful way,

when the Web link is further described by a relation type and a set of target attributes. A hypermedia control is the machine-readable description of *how* to activate an affordance. Hypermedia controls usually originate from a Web server and are discovered in-band while a Web client is interacting with the server. This way, Web servers can drive clients through Web applications dynamically, by taking their current state and other factors such as authorization into account. This is opposed to out-of-band interface descriptions that need to be preinstalled or hardcoded into clients (e.g., RPC, WS-\* Web services, HTTP services with fixed URI-method-response definitions).

W3C WoT makes use of two kinds of hypermedia controls: *Web links* [RFC8288], the well-established control to navigate the Web, and *Web forms* as a more powerful control to enable any kind of operation. Links are already used in other IoT standards and IoT platforms such as CoRE Link Format [RFC6690], OMA LWM2M [LWM2M], and OCF [OCF]. *Form* is a new concept that besides W3C WoT is also introduced by the *Constrained RESTful Application Language (CoRAL)* [CoRAL] defined by the IETF.

## § 6.6.1 Links

*Links* enable Consumers (or Web clients in the broader sense) to change the current context (cf. the set of resource representations currently rendered in the Web browser) or to include additional resources into the current context, depending on the relation between context and link target. Consumers do so by *dereferencing* the target URI, that is, fetching the resource representation by following a link.

W3C WoT follows the definitions of Web Linking [RFC8288], where a link is comprised of:

- a link context,
- a relation type,
- a link target, and
- optionally target attributes.

Link relation types are either a set of predefined tokens that are registered with IANA [IANA-RELATIONS], (e.g. stylesheet), or extension types in the form of URIs [RFC3986]. Extension relation types *MUST* be compared as strings using ASCII case-insensitive comparison, (c.f. ASCII case insensitive). (If they are serialized in a different format they are to be converted to URIs). Nevertheless, all-lowercase URIs *SHOULD* be used for extension relation types [RFC8288].

In the Web of Things, *links* are used for Discovery and to express relations between Things (e.g., hierarchical or functional) and relations to other documents on the Web (e.g., documentation manuals or alternative representations such as CAD models).



## § 6.6.2 Forms

Forms enable Consumers (or Web clients in the broader sense) to perform operations that go beyond dereferencing a URI (e.g., to manipulate the state of a Thing). Consumers do so by *filling out* and *submitting* the form to its submission target. This usually requires more detailed information about the contents of the (request) message than a link can provide (e.g., method, header fields, or other protocol options). Forms can be seen as a request template, where the provider pre-filled parts of the information according to its own interface and state, and left parts blank to be filled by the Consumers (or Web client in general).

~~W3C~~ WoT defines forms as new hypermedia control. Note that the definition in CoRAL is virtually identical, and hence compatible [CoRAL]. In CoRAL a form is comprised of:

- a form context,
- an operation type,
- a submission target,
- a request method, and
- optionally form fields.

A form can be viewed as a statement of "To perform an operation type operation on form context , issue a request method request to submission target " where the optional form fields may further describe the required request.

Form contexts and submission targets *MUST* both be Internationalized Resource Identifiers (IRIs) [RFC3987]. However, in the common case, they will also be URIs [RFC3986], because many protocols (such as HTTP) do not support IRIs.

Form context and submission target *MAY* point to the same resource or different resources, where the submission target resource implements the operation for the context.

The operation type identifies the semantics of the operation. Operation types are denoted similar to link relation types.

The request method *MUST* identify one method of the standard set of the protocol identified by the submission target URI scheme.

Form fields are optional and *MAY* further specify the expected request message for the given operation. Note that this is not limited to the payload, but may affect also protocol headers. Form fields *MAY* depend on the protocol used for the submission target as specified in the URI scheme. Examples are HTTP header fields, CoAP options, the protocol-independent media type [RFC2046]

including parameters (i.e., full content type) for the request payload, or information about the expected response.

#### NOTE

As of this specification, the well-known operation types are a fixed set that results from the WoT Interaction Model. Other specifications may define further well-known operation types that are valid for their respective document format or form serialization. Later versions of this specification or another specification may set up an IANA registry in the future to enable extension and a more generic Web form model that may be applied beyond WoT specifications.

## § 6.7 Protocol Bindings

A Protocol Binding is the mapping from an Interaction Affordance to concrete messages of a specific protocol such as HTTP [[RFC7231](#)], CoAP [[RFC7252](#)], or MQTT [[MQTT](#)]. It informs the Consumer *how* to activate the Interaction Affordance through a network-facing interface. The Protocol Bindings follow the Uniform Interface constraint of REST [[REST](#)] to support interoperability. Thus, not all communication protocols are eligible to implement Protocol Bindings for ~~W3C~~ WoT; the requirements are given in the assertions below.

In the door example given in [6.2 Affordances](#), the Protocol Binding corresponds to the door handle at the level of knob vs lever, which suggests *how* the door can be opened.

### § 6.7.1 Hypermedia-driven

Interaction Affordances *MUST* include one or more Protocol Bindings. Protocol Bindings *MUST* be serialized as [hypermedia controls](#) to be self-descriptive on how to activate the Interaction Affordance. The authority of the [hypermedia controls](#) can be the Thing itself, producing the TD document at runtime (based on its current state and including network parameters such as its IP address) or serving it from memory with only the current network parameters inserted. The authority can also be an external entity that has full and up-to-date knowledge of the Thing including its network parameters and internal structure (e.g., software stack). This enables a loose coupling between Things and Consumers, allowing for an independent lifecycle and evolution. The [hypermedia controls](#) *MAY* be cached outside the Thing and used for offline processing if caching metadata is available to determine the freshness.

### § 6.7.2 URIs

Hypermedia controls rely on URIs [RFC3986] to identify link and submission targets. Thereby, the URI scheme (the first component up to ":",) identifies the communication protocol to be used for Interaction Affordances with the Thing. W3C WoT refers to these protocols as transport protocols.

## § 6.8 Media Types

All data (a.k.a. content) exchanged when activating Interaction Affordances *MUST* be identified by a media type [RFC2046] in the Protocol Binding. Media types are labels to identify representation formats, for instance `application/json` for JSON [RFC8259] or `application/cbor` for CBOR [RFC7049]. They are managed by IANA.

Some media types might need additional parameters to fully specify the representation format used. Examples are `text/plain; charset=utf-8` or `application/ld+json; profile="http://www.w3.org/ns/json-ld#compacted"`. This needs to be considered in particular when describing data to be sent to Things. There might also be standardized transformations on the data such as content coding [RFC7231]. Protocol Bindings *MAY* have additional information that specifies representation formats in more detail than the media type alone.

Note that many media types only identify a generic serialization format that does not provide further semantics for its elements (e.g., XML, JSON, CBOR). Thus, the Interaction Affordance for structured data types *SHOULD* be associated with a *data schema* to provide more detailed syntactic metadata for the data exchanged. Details are further described in the WoT Thing Description specification [WOT-THING-DESCRIPTION].

## § 6.9 Internationalization


Web of Things support interoperable internationalization and allows to work with multilingual data such as for User Interfaces. The design and implementation of multilingual Web of Things implementations is guided by the Thing Description [WOT-THING-DESCRIPTION] specification. It describes how human-readable text in different languages can be applied based on established standards such as from [JSON-LD] and [BCP47].

## § 6.10 WoT System Components and their Interconnectivity

Section 6.1 [Fundamental Concepts](#) described the WoT architecture in terms of the abstract WoT architecture components such as [Things](#), [Consumers](#) and [Intermediaries](#). When those components are implemented as a software stack to take a specific role in the WoT architecture, such software stacks are called [Servients](#). Systems that are based on the WoT architecture involve [Servients](#), which are communicating with each other to achieve the goals of a system.

This section uses system configuration diagrams to illustrate how [Servients](#) work together to build systems based on the WoT architecture.


A [Thing](#) can be implemented by a [Servient](#). In a [Thing](#), a [Servient](#) software stack contains a representation of a [Thing](#) called [Exposed Thing](#), and makes its [WoT Interface](#) available to [Consumers](#) of the [Thing](#). This [Exposed Thing](#) may be used by other software components on the [Servient](#) (e.g., applications) to implement the behavior of the thing.



*Figure 24 Servient as a Thing*

On the other hand, [Consumers](#) are always implemented by [Servients](#), as they must be able to process the [Thing Description](#) (TD) format and must have a protocol stack that can be configured through [Protocol Binding](#) information contained in the TDs.

In a [Consumer](#), a [Servient](#) software stack provides a representation of a [Thing](#) called [Consumed Thing](#), and makes it available to those applications running on the [Servient](#) that need to process TDs to interact with [Things](#).



*Figure 25 Servient as a Consumer*

A [Consumed Thing](#) instance in the [Servient](#) software stack serves to separate the protocol level complexity from applications. It is communicating with [Exposed Things](#) on behalf of the application.

Similarly, an [Intermediary](#) is yet another WoT architecture component implemented by a [Servient](#). An [Intermediary](#) is located between a [Thing](#) and its [Consumers](#), performing the roles of both a

Consumer (to the Thing) and a Thing (to the Consumers). In an Intermediary, a Servient software stack contains the representations of both a Consumer (Consumed Thing) and a Thing (Exposed Thing).




Figure 26 Servient as an Intermediary

### § 6.10.1 Direct Communication

Figure 27 shows direct communication between a Thing, which is exposing Interaction Affordances through Thing Descriptions, and a Consumer that uses the Thing by means of the Interaction Affordances. Direct communication applies when both Servients use the same network protocol(s) and are accessible to each other.




Figure 27 High-level architecture of Consumer and Thing

An Exposed Thing is the software representation of a Thing abstraction, serving a WoT Interface of the Interaction Affordances provided by the Thing.

A Consumed Thing is the software representation of a remote Thing being consumed by a Consumer, serving as the interface to the remote Thing for the applications. A Consumer can generate a Consumed Thing instance by parsing and processing a TD document. Interactions between a Consumer and a Thing are performed by the Consumed Thing and the Exposed Thing exchanging messages over a direct network connection between them.

### § 6.10.2 Indirect Communication

In [Figure 28](#), a [Consumer](#) and a [Thing](#) connect to each other via an [Intermediary](#). An [Intermediary](#) is required if the [Servients](#) use different protocols or if they are on different networks that require authentication and provide access control (e.g. firewalls).



*Figure 28 High-level architecture with Intermediary*

An [Intermediary](#) combines [Exposed Thing](#) and [Consumed Thing](#) functionality. The functionality of [Intermediaries](#) includes relaying messages for the Interaction Affordances between a [Consumer](#) and a [Thing](#), optionally caching the [Thing](#)'s data for faster response, and transforming communication when the functionality of the [Thing](#) is extended by the [Intermediary](#). In an [Intermediary](#), a [Consumed Thing](#) creates a proxy object of the [Exposed Thing](#) of a [Thing](#), and a [Consumer](#) can access the proxy object (i.e., the [Exposed Thing](#) of the [Intermediary](#)) through its own [Consumed Thing](#).

[Consumer](#) and [Intermediary](#) can communicate in a different protocol than [Intermediary](#) and [Thing](#). For example, an [Intermediary](#) can provide a bridge between a [Thing](#) that uses CoAP and the application of a [Consumer](#) that uses HTTP.

Even when there are multiple different protocols used between [Intermediary](#) and [Things](#), [Consumer](#) can indirectly communicate with those [Things](#) using a single protocol through the [Intermediary](#). The same is true for the authentication. The [Consumed Thing](#) of a [Consumer](#) only needs to authenticate with the [Exposed Things](#) of the [Intermediary](#) using a single security mechanism, while the [Intermediary](#) might need multiple security mechanism to authenticate with different [Things](#).

Usually, an [Intermediary](#) generates the [Thing Description](#) for its proxy object based on the [Thing Description](#) of the originating [Thing](#). Depending on the requirements of the use cases, the TD for the proxy object may either use the same identifier as the TD of the original [Thing](#), or it gets assigned a new identifier. If necessary, a TD generated by an [Intermediary](#) *MAY* contain interfaces for other communication protocols.

## § 7. WoT Building Blocks

*This section is non-normative.*

The Web of Things (WoT) building blocks allow the implementation of systems that conform with the abstract WoT Architecture. The specifics of these building blocks are defined in separate specifications; this section provides an overview and a summary.

The WoT building blocks support each of the architectural aspects of a [Thing](#) discussed in 6.3 [Web Thing](#) and depicted in [Figure 20](#). The individual building blocks are shown in the context of an abstract [Thing](#) in [Figure 29](#). This is an abstract view and does not represent any particular implementation; instead it illustrates the relationship between the building blocks and the main architectural aspects of a [Thing](#). In this figure the WoT building blocks are highlighted with black outlines. The [WoT Thing Description](#) is a key building block that provides metadata describing a [Thing](#) and its network interface. Security, a cross-cutting concern, is separated into public and protected private components. The [WoT Scripting API](#) is optional and the [Binding Templates](#) are informative. The [WoT Discovery](#) building block defines mechanisms for distributing [Thing Descriptions](#); a [Thing](#) can provide [Thing Descriptions](#) directly, or they can be provided by a [Thing Description Directory](#) service.

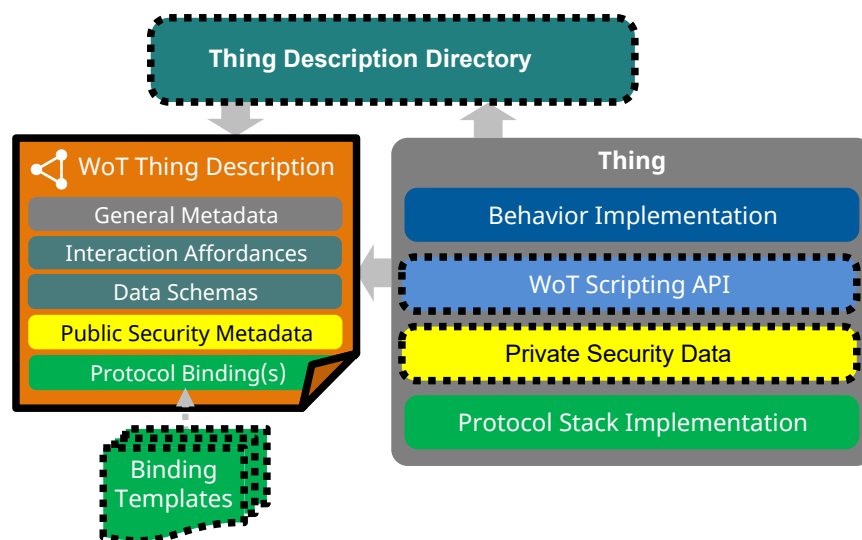


Figure 29 Relationship of WoT Building Blocks to the Architectural Aspects of a Thing.

In the following sections we will provide additional information on each WoT building block: the [WoT Thing Description](#), the [WoT Discovery](#) mechanisms, the [WoT Binding Templates](#), and the [WoT Scripting API](#). Security, although it is a cross-cutting concern, can be considered an additional building block.

## § 7.1 WoT Thing Description

The WoT Thing Description (TD) specification [[WOT-THING-DESCRIPTION](#)] defines an *information model* based on a semantic vocabulary and a *serialized representation based on JSON*. TDs provide rich metadata for Things in a way that is both human-readable and machine-readable. Both the information model and the representation format of TDs are aligned with Linked Data [[LINKED-DATA](#)], so that besides raw JSON processing, implementations may choose to make use of JSON-LD [[JSON-LD11](#)] and graph databases to enable powerful semantic processing of the metadata.

A Thing Description describes Thing instances with general metadata such as name, ID, descriptions, and also can provide relation metadata through links to related Things or other documents. TDs also contain Interaction Affordance metadata based on the interaction model defined in [6.5 Interaction Model](#); Public Security Metadata; and communications metadata defining Protocol Bindings. The TD can be seen as the *index.html* for Things, as it provides the entry point to learn about the services and related resources provided by a Thing, both of which are described using hypermedia controls.

For semantic interoperability, TDs may make use of a domain-specific vocabulary, for which explicit extension points are provided. However, development of any particular domain-specific vocabulary is currently out-of-scope of the [W3C](#) WoT standardization activity.

Three examples of potentially useful external IoT vocabularies are SAREF [[SAREF](#)], Schema Extensions for IoT [[IOT-SCHEMA-ORG](#)], and the [W3C](#) Semantic Sensor Network ontology [[VOCAB-SSN](#)]. Use of such external vocabularies in TDs is optional. In the future additional domain-specific vocabularies may be developed and used with TDs.

Overall, the WoT Thing Description building block fosters interoperability in two ways: First, TDs enable machine-to-machine communication in the Web of Things. Second, TDs can serve as a common, uniform format for developers to document and retrieve all the details necessary to create applications that can access IoT devices and make use of their data.

TDs must be known or accessible to be useful to other systems and devices. The WoT Discovery building block, described in more detail below, defines mechanisms by which this may be accomplished for both self-describing devices and for situations (brownfield devices, constrained devices, devices with specialized protocols, sleeping devices, etc.) in which having a separate service providing the TD is more suitable.

## § 7.2 Thing Model

The Thing Model defines a template-based model for Thing Description instances. A Thing Model has no instance specific and only partial communication- and security-based information. This kind of information is supplemented by the creation of Thing Description instantiation.



A Thing Model mainly describes the Properties, Actions, and Events and common metadata which then are available in all instantiated Thing Descriptions. This paradigm can be compared with abstract class or interface definitions (~Thing Model) in object-oriented programming to create objects (~Thing Descriptions). Such Thing Models are relevant for, e.g., mass production of IoT devices, onboarding scenarios as in cloud services, or to simulate Devices that have not yet been developed.

## § 7.3 Profiles

The WoT Profile Specification [[WOT-PROFILE](#)] defines Profiles, that enable *out of the box interoperability* among things and devices. Out of the box interoperability implies, that devices can be integrated into various application scenarios without deep level adaptations. Typically only minor configuration operations are necessary (such as entering a network key, or IP address) to use the device in a certain scenario. These actions can be done by anyone without specific training.

The HTTP Baseline Profile defines a mapping to HTTP(S). The HTTP Baseline Profile is complemented by two other profiles that provide protocol bindings for the event mechanism:

- HTTP SSE Profile for Server Sent Events
- HTTP Webhooks Profile for Webhooks

These profiles can be used stand alone, or a device can implement a combination of them. It is envisioned, that other profiles will be defined in the future, that contain mappings to other protocols.

### § 7.3.1 Profiling Methodology

A profile is a set of constraints and rules, which provide additional semantic guarantees that are applied to the WoT Thing Description specification. These constraints define a subset of valid WoT Thing Descriptions by defining additional rules on various aspects of the WoT Thing Description.

<i>Constraints on</i>	<i>Rationale</i>	<i>Examples</i>
vocabulary of Thing Description classes	guaranteed set of metadata fields	Make specific vocabulary terms mandatory, remove others
class relationships	unambiguous structure	limited cardinality, e.g. only one form per operation per interaction affordance.

values of vocabulary terms	simplified processing	Limit the length of characters per string. Always use arrays, where the spec permits a string or an array of strings.
data schemas	simplified processing	No arbitrary nested objects or arrays of arrays.
security	reduced implementation effort	Only a restricted set of security mechanisms.
protocol binding	guaranteed protocol semantics	limited protocol(s) and protocol features, Example: predefined mapping of http verbs (GET/PUT) to operation verbs, similar constraints for other protocols.

These constraints and rules fall into two categories:

- constraints on the data model.
- constraints on the protocol binding.

These two categories are orthogonal to each other, i.e. an information model that conforms to a profile can be mapped to different protocols. The protocol binding for each protocol may contain additional (protocol-specific) constraints.

A profile is not exclusive, i.e. a thing may conform to multiple profiles. Profiles can build on top of each other or overlap, extended profiles can be built on top of a baseline profile.

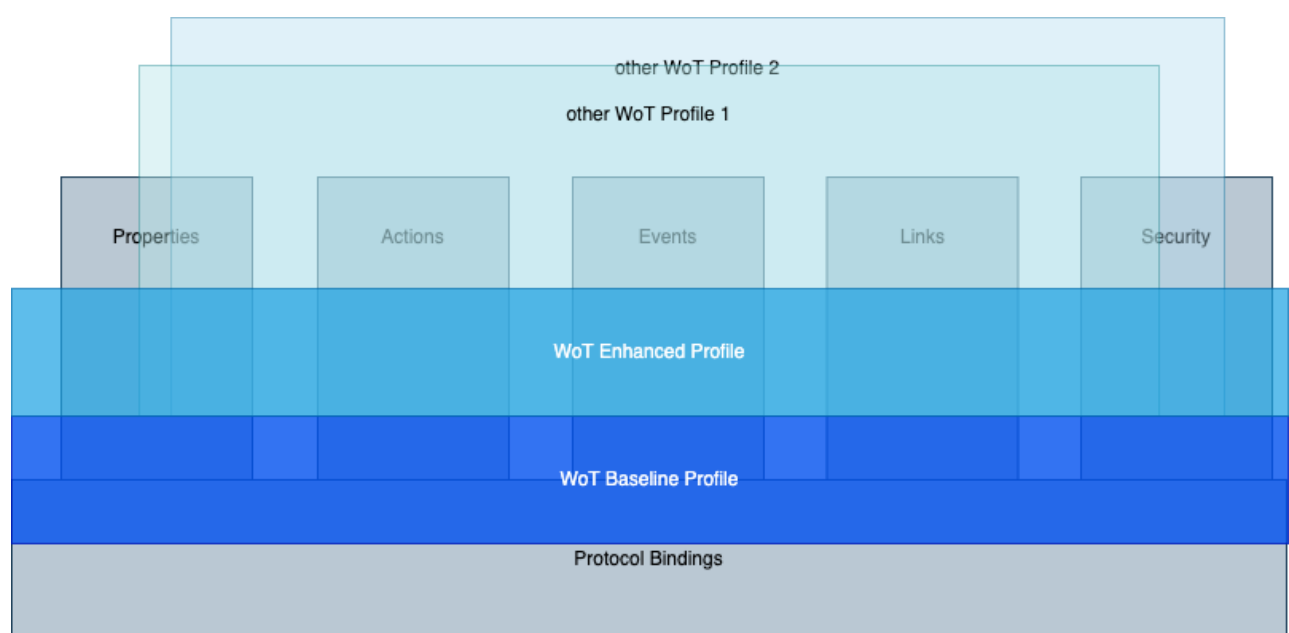


Figure 30 WoT Baseline Profile - Other Profiles

## § 7.4 WoT Discovery

The WoT Discovery specification [[WOT-DISCOVERY](#)] defines mechanisms for distributing and accessing WoT Thing Descriptions over the network. These mechanisms are meant to simplify access to Things and services and support their integration. WoT Discovery mechanisms are not limited by the bounds of local area networks and can also support remote discovery. They are also designed to act as extensions of existing search and discovery standards.

One of the primary design requirements of the WoT Discovery building block is to protect the security and privacy of both consumers and providers of WoT data and metadata. In particular, WoT Discovery mechanisms are designed not only to protect and control access to WoT Thing Descriptions, but also are designed to protect queries about such data and metadata, and to avoid accidental distribution of direct or inferrable metadata.

In order to accommodate existing discovery technologies while providing strong privacy and security, WoT Discovery uses a two-stage process. In the first stage, an "introduction" is made using one of several first-contact mechanisms. In the second stage, called "exploration", actual TDs are made available, but only after suitable authentication and authorization.

Two exploration mechanisms are provided, supporting both distribution of single TDs and directory services for multiple TDs:

- An ordinary web resource provided by an HTTP server can be used to distribute a single TD. This can also be used for self-description, to allow a Thing to provide its own WoT Thing Description directly.
- Directory services provide searchable databases of WoT Thing Descriptions. Both lightweight syntactic query mechanisms and more powerful semantic query mechanisms are supported by directory services.

Use of these mechanisms is encouraged, but not mandatory.

### § 7.4.1 Introduction Mechanisms

Introduction mechanisms are not intended to provide strong security or privacy guarantees, and this allows a variety of existing discovery mechanisms with relatively weak security to be used, as long as they satisfy the following requirements:

- Any introduction mechanism used for WoT Discovery must not provide metadata but instead should simply result in one or more opaque URLs at which metadata may be accessed, if and

only if the user can authenticate and has the appropriate authorization (access rights).

- [Introduction](#) URLs must not themselves include any metadata, e.g. a human-readable device type or name. Randomly generated URLs are recommended.

## § 7.4.2 Exploration Mechanisms

After a suitable authentication and authorization process (based on existing mechanisms to protect web services and existing protocol security negotiation mechanisms), [WoT Discovery](#) defines a set of second stage "[exploration](#)" mechanisms that provide actual access to metadata.

The first and simplest [exploration](#) mechanism defines how to fetch a [TD](#) directly from the device itself. This supports ad-hoc peer-to-peer discovery. However, in many circumstances (including but not limited to management of large collections of TDs), an alternative [exploration](#) mechanism, the [WoT Thing Description Directory \(TDD\)](#) service, is more appropriate.

A [TDD](#) provides sophisticated (but protected) query mechanisms to explore and retrieve WoT metadata. [TDDs](#) also provide change notification mechanisms to support secure distribution of [TD](#) updates to authorized consumers. Any [WoT Discovery](#) "exploration" implementation must only provide metadata and [TDs](#) after appropriate best-practice authentication and authorization. Suitable best-practice security mechanisms for authentication and authorization for different circumstances are discussed in [[WOT-SECURITY](#)]. Suitable mechanisms for managing access controls and keys are also discussed in [[SOLID](#)].

[TDDs](#) are not just a convenience feature but are essential in several WoT use cases. Using a [TDD](#) is appropriate when a [Thing](#) has resource restrictions (e.g., limited memory space, limited power), c.f. [Device Categories](#), uses a specialized protocol requiring translation by an intermediary (in which case the [TD](#) provided by the [TDD](#) would describe the translated network interface supported by the intermediary), or when it is necessary to use an existing device (often called "brownfield" devices) as part of the Web of Things but the device itself cannot be easily modified to self-host a [TD](#).

[TDs](#) can be registered with a [TDD](#) by the devices themselves or by an external agent. Use of an external agent will usually be required in the case of brownfield devices.

The [Discovery](#) mechanisms defined in [[WOT-DISCOVERY](#)] are not mandatory, but they are designed to meet the above requirements and their use will enhance interoperability and usability.

## § 7.5 WoT Binding Templates

The IoT uses a variety of protocols for accessing devices, since no single protocol is appropriate in all contexts. Thus, a central challenge for the Web of Things is to enable interactions with the plethora of different IoT platforms (e.g., OMA LWM2M, OPC UA, oneM2M) and devices that do not follow any particular standard, but provide an eligible interface over a suitable network protocol. WoT is tackling this variety through Protocol Bindings, which must meet a number of constraints (see [6.7 Protocol Bindings](#)).

The Binding Templates address the aspect that an application client (a Consumer) can use the Thing Description to extract the specific metadata of the protocols (e.g. about HTTP, MQTT, Modbus, etc.), payload formats (binary, JSON, CBOR, EXI, etc.) and their usage in an IoT platform context (e.g., ECHONET, OPC UA). The metadata can be passed to a network implementation interface to establish interaction with the Thing and to serialize / deserialize the payload messages. In that context, Binding Templates include three kinds of bindings: Protocol bindings, Payload format bindings, and Platform bindings.

The non-normative WoT Binding Templates specification [[WOT-BINDING-TEMPLATES](#)] provides a collection of blueprints that give guidance on how to interact with different Things that use different transport protocols, content types or that are different IoT platforms or standards (e.g., OCF, oneM2M, OMA LWM2M, OPC UA) which use certain combinations of transport protocols and content types. When describing a particular IoT device or platform, the corresponding Binding Template can be used to look up the communication metadata that is to be provided in the Thing Description to support that platform.

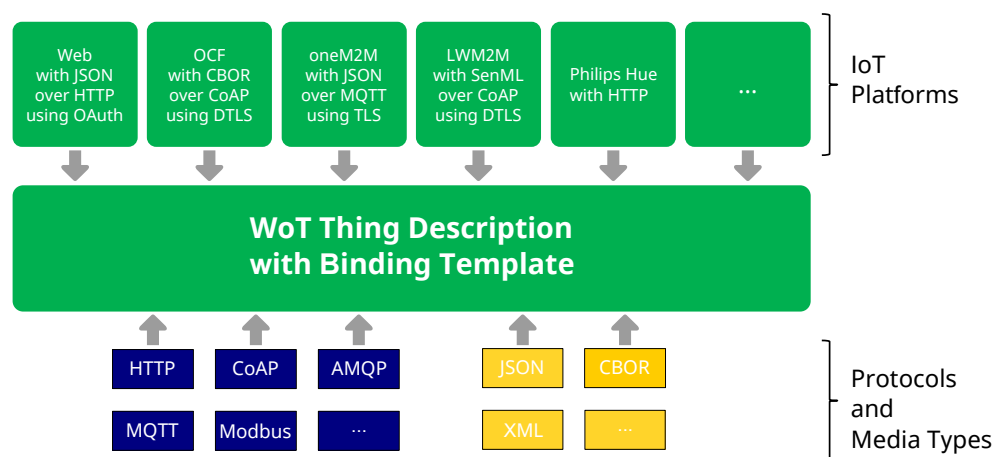


Figure 31 From Binding Templates to Protocol Bindings

[Figure 31](#) shows how Binding Templates are used. Based on the protocol, media type or platform binding, a TD is created. The Consumer that is processing a TD implements the required Binding Template that is present in the TD by including a corresponding protocol stack, media type encoder/decoder or platform stack and by configuring the stack (or its messages) according to the information given in the TD such as serialization format of the messages and header options.

The Binding Templates spans three dimensions:

- **Protocol**

The Web of Things uses the term transport protocol for the underlying, standardized application-layer protocol without application-specific options or subprotocol mechanisms. The URI scheme found in the form of an Interaction Affordance and the associated protocol endpoint URI contains the information required to identify the transport protocol, e.g., HTTP, CoAPS, or WebSocket through http, coaps, or ws, respectively.

- **Subprotocol**

Transport protocols may have to use information provided by an extension mechanisms to interact successfully. Such subprotocols cannot be identified from the URI scheme alone and must be declared explicitly using the subprotocol field inside forms. Examples for HTTP are long polling [[RFC6202](#)] (longpoll) or Server-Sent Events [[HTML](#)] (sse).

- **Payload Representation**

Representation formats (a.k.a. serializations) used for exchanging data can differ between protocols, IoT Platforms and standards. The media type [[RFC2046](#)] identifies these formats, while media type parameters may specify them further. Forms may contain the media type and optional parameters in additional form fields such as a content type field known from HTTP, which combines media type and its potential parameters (e.g., text/plain; charset=utf-8).

- **Platforms and Standards**

IoT Platforms and standards often introduce specific modifications at the application layer such as platform-specific HTTP header fields or CoAP options. They can also contain specific payload structures on top of the media types that need to be accurately incorporated into each message between the Thing and its Consumer. Both the protocol modification and payload representation specific to different platforms can be represented in TDs using the best practices documented in binding templates for platforms and standards.

## § 7.6 WoT Scripting API

The WoT Scripting API is an optional "convenience" building block of ~~W3C~~ WoT that eases IoT application development by providing an ECMAScript-based API [[ECMAScript](#)] similar to the Web browser APIs. By integrating a scripting runtime system into the WoT Runtime, the WoT

Scripting API enables using portable application scripts that define the behavior of Things, Consumers, and Intermediaries.

Traditionally, IoT device logic is implemented in firmware, which results in productivity constraints similar to that of embedded development, including a relatively complex update process. The WoT Scripting API in contrast enables implementing device logic by reusable scripts executed in a runtime system for IoT applications. It is similar to that of a Web browser, and aims to improve productivity and reduce integration costs. Furthermore, standardized APIs enable portability for application modules, for instance, to move compute-intense logic from a device up to a local gateway, or to move time-critical logic from the cloud down to a gateway or edge node.

The non-normative WoT Scripting API specification [[WOT-SCRIPTING-API](#)] defines the structure and algorithms of the programming interface that allows scripts to discover, consume, and expose WoT Thing Descriptions. The runtime system of the WoT Scripting API instantiates local objects that act as an interface to other Things and their Interaction Affordances (Properties, Actions, and Events). It also allows scripts to expose Things, that is, to define and implement Interaction Affordances and publish a Thing Description.

## § 7.7 WoT Security and Privacy Guidelines

Security is a cross-cutting concern and should be considered in all aspects of system design. In the WoT architecture, security is supported by certain explicit features, such as support for Public Security Metadata in TDs and by separation of concerns in the design of the WoT Scripting API. The specification for each building block also includes a discussion of particular security and privacy considerations of that building block. Another non-normative specification, the *WoT Security and Privacy Guidelines* [[WOT-SECURITY](#)], provides additional cross-cutting security and privacy guidance.

## § 8. Abstract Servient Architecture

*This section is non-normative.*

As defined in [6.10 WoT System Components and their Interconnectivity](#), a Servient is a software stack that implements the WoT building blocks presented in the previous section. Servients can host and expose Things and/or consume Things (i.e., host Consumers). Depending on the Protocol Binding, Servients can perform in both server and client role, hence the portmanteau naming.

The previous section describes how the WoT building blocks conceptually relate to each other and how they correspond to the abstract WoT Architecture (see [6. Abstract WoT System Architecture](#)).

When implementing these concepts, a more detailed view is necessary that takes certain technical aspects into account. This section describes the detailed architecture of a Servient implementation.

[Figure 32](#) shows a Servient implementation that is using the (optional) WoT Scripting API building block. Here, the WoT Runtime is also a Scripting Runtime system that, in addition to managing the WoT-specific aspects, also interprets and executes the application scripts. Servients that support the WoT Scripting API usually run on powerful devices, edge nodes, or in the cloud (c.f. [Device Categories](#)). The WoT Architecture does not limit the application-facing API of the WoT Runtime to JavaScript/ECMAScript. Also other runtime systems can be used to implement a Servient.

Section 8.8.1 Native WoT API presents an alternative Servient implementation without the WoT Scripting API building block. The WoT Runtime may use any programming language for its application-facing API. Usually, it is the native language of the Servient software stack, for instance C/C++ for embedded Servients or Java for cloud-based Servients. It may also be an alternative scripting language such as Lua to combine the benefits of application scripts with low resource consumption.

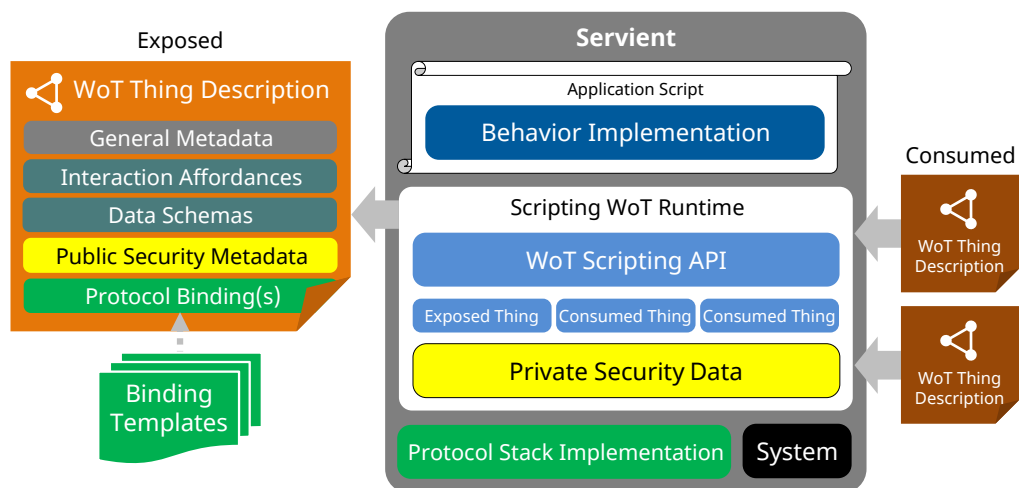


Figure 32 Implementation of a Servient using the WoT Scripting API

The role and functionality of each module shown in [Figure 32](#) is explained in the following sections.

## § 8.1 Behavior Implementation

The *behavior* defines the overall application logic of a Thing, which has several aspects:

It includes *autonomous behavior* of Things (e.g., sampling of sensors or control loops for actuators), the *handlers* for Interaction Affordances (i.e., the concrete actions taken when an



affordance is activated), *Consumer behavior* (e.g., controlling a Thing or realizing mashups), and *Intermediary behavior* (e.g., simply proxying a Thing or composing virtual entities). The behavior implementation within a Servient defines which Things, Consumers, and Intermediaries are hosted on this component.

[Figure 32](#) depicts Servients that are implementing the optional WoT Scripting API building block, where portable application scripts written in JavaScript [[ECMAScript](#)] define the behavior. They are executed by a scripting runtime system that is part of the WoT Runtime (when providing the WoT Scripting API or any other script-based API). They are portable, as they are written against the common WoT Scripting API definitions, and hence can be executed by any Servient featuring this building block. This makes it possible to shift application logic between system components, for instance moving a Consumer from the cloud to an edge node to meet networking requirements, or to move an Intermediary to the cloud to fulfill growing resource demands. Portable applications enable to 'install' additional behavior after the deployment of a Servient.

In principle, any programming language and API can be used in order to define the behavior of a Thing, as long as the Interaction Affordances are presented externally through a WoT Interface. The adaption between application-facing API and the protocol stack is handled by the WoT Runtime. See [8.8.1 Native WoT API](#) for behavior implementation without the WoT Scripting API building block.

## § 8.2 WoT Runtime

Technically, the Thing abstraction and its Interaction Model is implemented in a runtime system. This WoT Runtime maintains the execution environment for the behavior implementation and is able to expose and/or consume Things, and hence must be able to fetch, process, serialize, and serve WoT Thing Descriptions.

Every WoT Runtime has an application-facing interface (i.e., an API) for the behavior implementation. The optional WoT Scripting API building block shown in [Figure 32](#) defines such an application-facing interface that follows the Thing abstraction and enables the deployment of behavior implementations during runtime through application scripts. See [8.8.1 Native WoT API](#) for alternative APIs, which can also only be available during compile time. In general, application logic should be executed in sandboxed execution environments that prevent unauthorized access to the management aspects of the WoT Runtime from the application code, in particular the Private Security Data. In multi-tenant Servients, different tenants should also be prevented from accessing each other's data without authorization.

A WoT Runtime needs to provide certain operations to manage the lifecycle of Things, or more precisely their software abstractions and descriptions. A lifecycle management (LCM) system may encapsulate those lifecycle operations within a Servient and use internal interfaces to realize the

lifecycle management. The details of such operations vary among different implementations. The WoT Scripting API includes LCM functionality, and hence represents one possible implementation of such a system.

The WoT Runtime must interface with the protocol stack implementation of the Servient, as it decouples the behavior implementation from the details of the Protocol Bindings. The WoT Runtime usually also interfaces with the underlying system, for instance, to access local hardware such as attached sensors and actuators or to access system services such as storage. Both interfaces are implementation-specific, yet the WoT Runtime must provide the necessary adaption to the Thing abstraction it implements.

## § 8.3 WoT Scripting API

The WoT Scripting API building block defines an ECMAScript API that closely follows the WoT Thing Description specification [[WOT-THING-DESCRIPTION](#)]. It defines the interface between behavior implementations and a scripting-based WoT Runtime. Other, simpler APIs may be implemented on top of it, similar to, for instance, jQuery for the Web browser APIs.

See [[WOT-SCRIPTING-API](#)] for more details.

## § 8.4 Exposed Thing and Consumed Thing Abstractions

The WoT Runtime instantiates software representations of Things based on their TDs. These software representations provide the interface towards the behavior implementation.

The Exposed Thing abstraction represents a Thing hosted locally and accessible from the outside through the protocol stack implementation of the Servient. The behavior implementation can fully control Exposed Things by defining their metadata and Interaction Affordances, and providing their autonomous behavior.

The Consumed Thing abstraction represents a remotely hosted Thing for Consumers that needs to be accessed using a communication protocol. Consumed Things are proxy objects or stubs. The behavior implementation is restricted to reading their metadata and activating their Interaction Affordances as described in the corresponding TD. Consumed Things can also represent system features such as local hardware or devices behind proprietary or legacy communication protocols. In this case, the WoT Runtime must provide the necessary adaptation between system API and Consumed Thing. Furthermore, it must provide corresponding TDs and make them available to the behavior implementation, for instance, by extending whatever discovery mechanism is provided by

the WoT Runtime through the application-facing API (e.g., the `discover()` method defined in the WoT Scripting API [[WOT-SCRIPTING-API](#)]).

When using the WoT Scripting API, Exposed Thing and Consumed Thing are JavaScript objects, which can be created, operated on, and destroyed by application scripts. However, access may be restricted through a security mechanism, for instance, in multi-tenant Servients.

## § 8.5 Private Security Data

Private security data, such as a secret key for interacting with the Thing, is also conceptually managed by the WoT Runtime, but is intentionally not made directly accessible to the application. In fact, in the most secure hardware implementations, such Private Security Data is stored in a separate, isolated memory (e.g., on a secure processing element or TPM) and only an abstract set of operations (possibly even implemented by an isolated processor and software stack) is provided that limit the attack surface and prevent external disclosure of this data. Private Security Data is used transparently by the Protocol Binding to authorize and protect the integrity and confidentiality of interactions.

## § 8.6 Protocol Stack Implementation

The protocol stack of a Servient implements the WoT Interface of the Exposed Things and is used by Consumers to access the WoT Interface of remote Things (via Consumed Things). It produces the concrete protocol messages to interact over the network. Servients may implement multiple protocols, and hence support multiple Protocol Bindings to enable interaction with different IoT Platforms.

In many cases, where standard protocols are used, generic protocol stacks can be used to produce the platform-specific messages (e.g., one for HTTP(S) dialects, one for CoAP(S) dialects, and one for MQTT solutions, etc.). In this case, the communication metadata from the Thing Description is used to select and configure the right stack (e.g., HTTP with the right header fields or CoAP with the right options). Parsers and serializers for the expected payload representation format (JSON, CBOR, XML, etc.) as identified by the media type [[RFC2046](#)] can also be shared across these generic protocol stacks.

See [[WOT-BINDING-TEMPLATES](#)] for details.

## § 8.7 System API

An implementation of a WoT Runtime may provide local hardware or system services to behavior implementations through the Thing abstraction, as if they were accessible over a communication protocol. In this case, the WoT Runtime should enable the behavior implementation to instantiate Consumed Things that internally interface with the system instead of the protocol stack. This can be done by listing such system Things, which are only available in the local WoT Runtime, in the results of the discovery mechanism provided by the application-facing WoT Runtime API.

A device may also be physically external to a Servient, but connected via a proprietary protocol or a protocol not eligible as WoT Interface (see 6.7 [Protocol Bindings](#)). In this case, the WoT Runtime may access legacy devices with such protocols (e.g., ECHONET Lite, BACnet, X10, I2C, SPI, etc.) through proprietary APIs, but may again choose to expose them to the behavior implementation via a Thing abstraction. A Servient can then act as gateway to the legacy devices. This should only be done if the legacy device cannot be described directly using a WoT Thing Description.

The behavior implementation may also access local hardware or system services (e.g., storage) through a proprietary API or other means. This is, however, out of scope of the ~~W3C~~ WoT standardization, as it hinders portability.

## § 8.8 Alternative Servient and WoT Implementations

The WoT Scripting API building block is optional. Alternative Servient implementations are possible, where the WoT Runtime offers an alternative API for the application logic, which may be written in any programming language.

Furthermore, devices or services unaware of ~~W3C~~ WoT can still be consumed, when it is possible to provide a well-formed WoT Thing Description for them. In this case, the TD describes a WoT Interface of a Thing that has a black-box implementation.

### § 8.8.1 Native WoT API

There are various reasons why a developer may choose to implement a Servient without using the WoT Scripting API. This may be due to insufficient memory or computing resources, so the developer cannot use the required software stack or a fully-featured scripting engine. Alternatively, to support their use case (for example, a proprietary communications protocol) the developer may have to use specific functions or libraries only available through a particular programming environment or language.

In this case, a [WoT Runtime](#) can still be used, but with an equivalent abstraction and functionality exposed using an alternative application-facing interface instead of the [WoT Scripting API](#). Except for the latter, all block descriptions in [8. Abstract Servient Architecture](#) are also valid for [Figure 33](#).

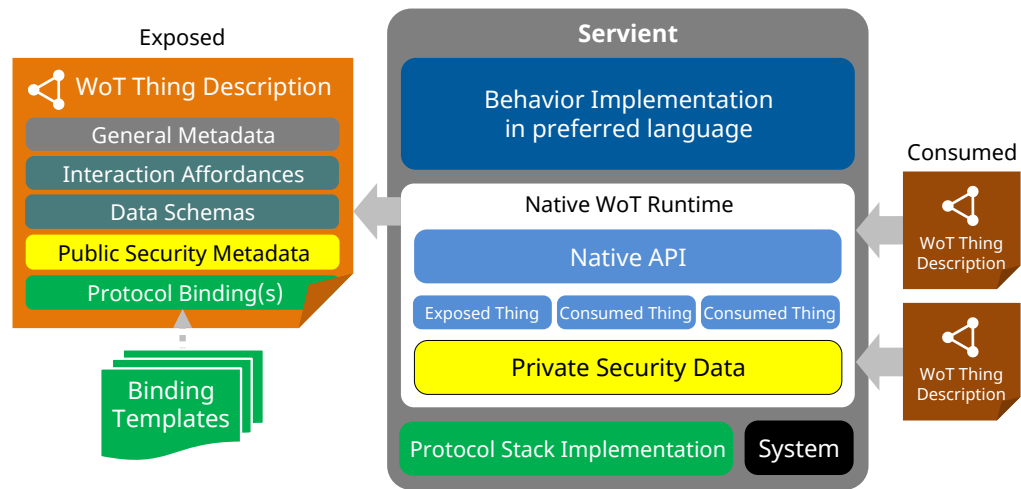


Figure 33 Implementation of a Servient Using a Native WoT API

## 8.8.2 Thing Description for Existing Devices

It is also possible to integrate *existing* IoT devices or services into the W3C Web of Things and to use them as Things by creating a Thing Description for these devices or services. Such a TD can either be created manually or via a tool or service. For example, a TD could be generated by a service that provides automatic translation of metadata provided by another, ecosystem-dependent machine-readable format. This can only be done, however, if the target device is using protocols that can be described using a Protocol Binding. The requirements for this are given in [6.7 Protocol Bindings](#). Much of the previous discussion also implies that a Thing provides its own Thing Description. While this is a useful pattern it is not mandatory. In particular, it may not be possible to modify existing devices to provide their own Thing Description directly. In this case the Thing Description will have to be provided separately using a service such as a directory or some other external and separate distribution mechanism.

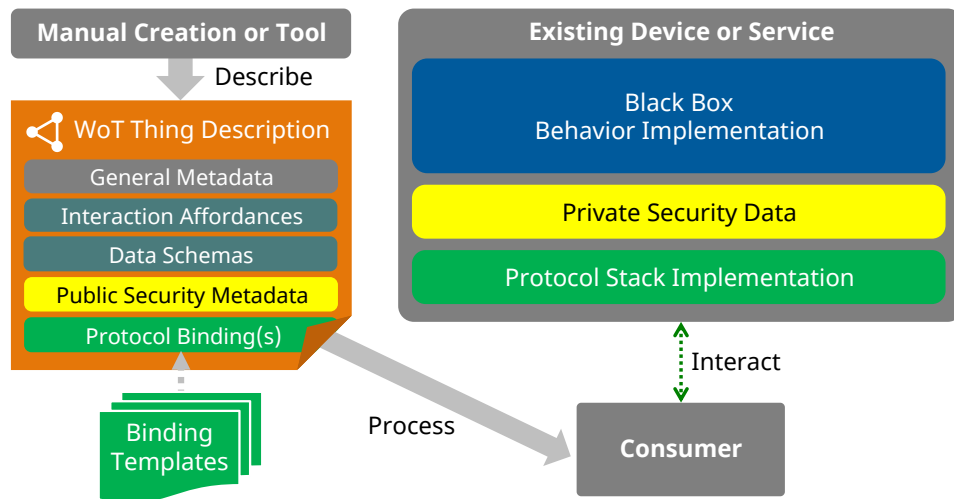


Figure 34 Integration of Existing IoT Devices into W3C WoT

## § 9. Example WoT Deployments

*This section is non-normative.*

This section provides various examples of how the Web of Things (WoT) abstract architecture may be instantiated when devices and services that implement the Thing and Consumer roles are connected in various concrete topologies and deployment scenarios. These topologies and scenarios are not normative, but are enabled and supported by the WoT abstract architecture.

Before discussing specific topologies, we will first review the roles that Things and Consumers can play in a WoT network and the relationships they have with the Exposed Thing and Consumed Thing software abstractions. Exposed Thing and Consumed Thing are internally available to the behavior implementations of Servients in the roles of Things and Consumers, respectively.

### § 9.1 Thing and Consumer Roles

A Servient in the role of a Thing creates an Exposed Thing based on a Thing Description (TD). TDs are published and made available to other Servients that are in the roles of Consumers or Intermediaries. TDs may be published in various different ways: the TD might be registered with a management system such as a Thing Description Directory service, or a Thing may provide the requesters with a TD upon receiving a request for a TD. It is even possible to statically associate a TD with Thing in certain application scenarios.

A Servient in the role of a Consumer obtains the TD of a Thing using a discovery mechanism and creates a Consumed Thing based on the obtained TD. The concrete discovery mechanism depends on the individual deployment scenario: It could be provided by a management system such as a Thing Description Directory, a discovery protocol, through static assignment, etc. As discussed elsewhere, use of the Discovery mechanisms defined in [WOT-DISCOVERY] is recommended whenever possible.

However, it should be noted that TDs describing devices associated with an identifiable person may potentially be used to infer privacy-sensitive information. Constraints on the distribution of such TDs must therefore be incorporated into any concrete TD discovery mechanism. If possible, steps to limit the information exposed in a TD may also have to be taken, such as filtering out IDs or human-readable information if this is not strictly necessary for a particular use case. Privacy issues are discussed at a high level in 11. Privacy Considerations and a more detailed discussion is given in the [WOT-THING-DESCRIPTION] specification and are addressed in the [WOT-DISCOVERY] specification.

Internal system functions of a device, such as interacting with attached sensors and actuators, can also optionally be represented as Consumed Thing abstractions.

The functions supported by the Consumed Thing are provided to the Consumer's behavior implementation through a programming language interface. In the WoT Scripting API, Consumed Things are represented by objects. The behavior implementation (that is, the application logic) running in a Thing can engage through Interaction Affordances with Consumers by using the programming language interface provided by the Exposed Thing.

A Thing does not necessarily represent a physical device. Things can also represent a collection of devices, or virtual services running in a gateway or in the cloud. Likewise, a Consumer may represent an application or service running on a gateway or cloud. Consumers can also be implemented on edge devices. In Intermediaries, a single Servient performs both the roles of a Thing and a Consumer simultaneously which are sharing a single WoT Runtime.

## § 9.2 Topology of WoT Systems and Deployment Scenarios

Various topologies and deployment scenarios of WoT systems are discussed in this section. These are only example patterns and other interconnection topologies are also possible. The topologies described here are derived from the Web of Things Use Cases and Requirements [WOT-USE-CASES-REQUIREMENTS].

### § 9.2.1 Consumer and Thing on the Same Network

In the simplest interconnection topology, illustrated by [Figure 35](#), the Consumer and Thing are on the same network and can communicate directly with each other without any intermediaries. One use case where this topology arises is when the Consumer is an orchestration service or some other IoT application running on a gateway and the Thing is a device interfacing to a sensor or an actuator. However, the client/server relationship could be easily reversed; the client could be a device in the Consumer role accessing a service running as a Thing on a gateway or in the cloud.



*Figure 35 Consumer and Thing on the Same Network*

If the Thing is in the cloud and the Consumer is on a local network (see [Figure 1](#) for an example in a Smart Home use case) the actual network topology may be more complex, for example requiring NAT traversal and disallowing certain forms of discovery. In such cases one of the more complex topologies discussed later may be more appropriate.

## § 9.2.2 Consumer and Thing Connected via Intermediaries

An Intermediary plays both Thing and Consumer roles on the network and supports both the Exposed Thing and Consumed Thing software abstractions within its WoT Runtime. Intermediaries can be used for proxying between devices and networks or for Digital Twins.

### § 9.2.2.1 Intermediary Acting as a Proxy

One simple application of an Intermediary is a proxy for a Thing. When the Intermediary acts as a proxy, it has interfaces with two separate networks or protocols. This may involve the implementation of additional security mechanisms such as providing TLS endpoints. Generally proxies do not modify the set of interactions, so the TD exposed by the Intermediary will have the same interactions as the consumed TD, however the connection metadata is modified.

To implement this proxy pattern, the Intermediary obtains a TD of a Thing and creates a Consumed Thing. It creates a proxy object of the Thing as a software implementation that has the same Interaction Affordances. It then creates a TD for the proxy object with a new identifier and possibly with new communications metadata (Protocol Bindings) and/or new Public Security



Metadata. Finally, an Exposed Thing is created based on this TD, and the Intermediary notifies other Consumers or Intermediaries of the TD via an appropriate publication mechanism.



*Figure 36 Consumer and Thing Connect via an Intermediary Acting as a Proxy*

#### § 9.2.2.2 Intermediary Acting as a Digital Twin

More complex Intermediaries may be known as Digital Twins. A Digital Twin may or may not modify the protocols or translate between networks, but they provide additional services, such as state caching, deferred updates, or even predictive simulation of the behavior of the target device. For example, if an IoT device has limited power, it may choose to wake up relatively infrequently, synchronize with a Digital Twin, and then immediately go to sleep again. In this case, typically the Digital Twins runs on a less power-constrained device (such as in the cloud or on a gateway), (c.f. Device Categories) and is able to respond to interactions on the constrained device's behalf. Requests for the current state of properties may also be satisfied by the Digital Twins using cached state. Requests that arrive when the target IoT device is sleeping may be queued and sent to it when it wakes up. To implement this pattern, the Intermediary, i.e., the Digital Twin needs to know when the device is awake. The device implementation as a Thing may need to include a notification mechanism for that. This could be implemented using a separate Consumer/Thing pair, or by using Event interactions for this purpose.

#### § 9.2.3 Devices in a Local Network Controlled from a Cloud Service


In Smart Home use cases, devices (sensors and home appliances) connected to a home network are often monitored and, in some cases, also controlled by cloud services. There is usually a NAT device between the home network to which the devices are connected and the cloud. The NAT device translates IP addresses as well as often providing firewall services, which block connections selectively. The local devices and cloud services can only communicate with each other if the communication can successfully traverse the gateway.

A typical structure, adopted in ITU-T Recommendation Y.4409/Y.2070 [Y.4409-Y.2070], is shown in [Figure 37](#). In this structure there is both a local and a remote Intermediary. The local Intermediary aggregates the Interaction Affordances from multiple Thing into a (set of) Exposed

Things, which can all be mapped onto a common protocol (for example, HTTP, with all interactions mapped to a single URL namespace with a common base server and using a single port). This provides the remote Intermediary with a simple way to access all the Things behind the NAT device, assuming the local Intermediary has used a converged protocol that can traverse the NAT device and has some way to expose this service to the Internet (STUN, TURN, dynamic DNS, etc.). In addition, the local Intermediary can function as a Thing proxy, so even when the connected Things each use a different protocol (HTTP, MQTT, CoAP, etc.) and/or a different set of ecosystem conventions, the Exposed Thing can converge them into a single protocol so that Consumers do not need to be aware of the various protocols the Things use.

In [Figure 37](#), there are two clients connected to the remote Intermediary, which has aggregated the services that reside behind the NAT border and may provide additional protocol translation or security services. In particular, the local Intermediary may be on a network with limited capacity and making that service directly available to all users may not be feasible. In this case access to the local Intermediary is *only* provided to the remote Intermediary. The remote Intermediary then implements a more general access control mechanism and may also perform caching or throttling to protect the consumer from excess traffic. Those consumers also will use a single protocol suitable for the open Internet (e.g., HTTPS) to communicate with the Intermediary, which makes the development of clients much simpler.

In this topology there is NAT and firewall functionality between the consumers and things, but the local and remote Intermediaries work together to tunnel all communications through the firewall, so the consumers and things need to know nothing about the firewall. The paired Intermediaries also protect the home devices by providing access control and traffic management.



*Figure 37 Cloud Applications Implemented as Consumers Connected to Local Devices implemented as Things via Paired Intermediaries*

In more difficult cases the NAT and firewall traversal may not work exactly as shown. In particular, an ISP may not support publicly accessible addresses, or STUN/TURN and/or dynamic DNS may not be supported or available. In this case the Intermediaries may alternative reverse the client/server roles between them to set up an initial connection (with the local Intermediary first connecting to the remote Intermediary in the cloud), then the pair of Intermediaries may establish a

tunnel (using for example, a Secure WebSocket, which uses TLS to protect the connection). The tunnel can then be used to encode all communications between the Intermediaries using a custom protocol. In this case the initial connection can still be made over HTTPS using standard ports, and from the local Intermediary to the remote Intermediary identically to a normal browser/web server interaction. This should be able to traverse most home firewalls, and since the connection is outgoing, network address translation will not cause any problems. However, even though a custom tunneling protocol is needed, the remote Intermediary can still translate this custom protocol back into standard external protocols. The connected Consumers and Things do not need to know about it. It is also possible to extend this example to use cases where both Things and Consumers can connect on either side of the NAT boundary. This however also requires a bidirectional tunnel to be established between the two Intermediaries.

#### § 9.2.4 Discovery Using a Thing Description Directory

Once local devices (and possibly services) can be monitored or controlled by services on cloud, a variety of additional services can be built on top. For example, a cloud application could change a device's operating condition based on an analysis of collected data.

However when the remote Intermediary is a part of a cloud platform servicing client applications, the clients need to be able to find device information by, for example, accessing a directory of connected devices. For simplicity in the figure below we have assumed all local devices are implemented as Things and all cloud applications as Consumers. To make the metadata of local devices implemented as Things available to the cloud applications, their metadata can be registered with a Thing Description Directory service. This metadata is specifically the TDs of the local devices modified to reflect the Public Security Metadata and communication metadata (Protocol Bindings) provided by the remote Intermediary. A client application then can obtain the metadata it needs to communicate with local devices to achieve its functionality by querying the Thing Description Directory.




Figure 38 Cloud Service with Thing Description Directory

In more complex situations, not shown in the figure, there may also be cloud services that act as Things. These can also register themselves with a Thing Description Directory. Since a Thing Description Directory is a Web service, it should be visible to the local devices through the NAT or firewall device and its interface can even be provided with its own TD. Local devices acting as Consumers can then discover the Things in the cloud via a Thing Description Directory and connect to the Things directly or via the local Intermediary if, for instance, protocol translation is needed.

## § 9.2.5 Service-to-Service Connections Across Multiple Domains

Multiple cloud eco-systems each based on different IoT platforms can work together to make a larger, system-of-systems eco-system. Building on the previously discussed structure of a cloud application eco-system, the figure below shows two eco-systems connected to each other to make a system-of-systems. Consider the case in which a client in one eco-system (i.e., Consumer A below) needs to use a server in another eco-system (i.e., Thing B below). There is more than one mechanism to achieve this cross eco-systems application-device integration. Below, two mechanisms are explained, each using a figure, to show how this can be achieved.

### § 9.2.5.1 Connection Through Thing Description Directory Synchronization

In [Figure 39](#), two instances of a Thing Description Directory synchronize information, which makes it possible for Consumer A to obtain the information of Thing B through Thing Description Directory A. As described in previous sections, remote Intermediary B maintains a shadow implementation of Thing B. By obtaining the TD of this shadow device, Consumer A is able to use Thing B through the remote Intermediary B.

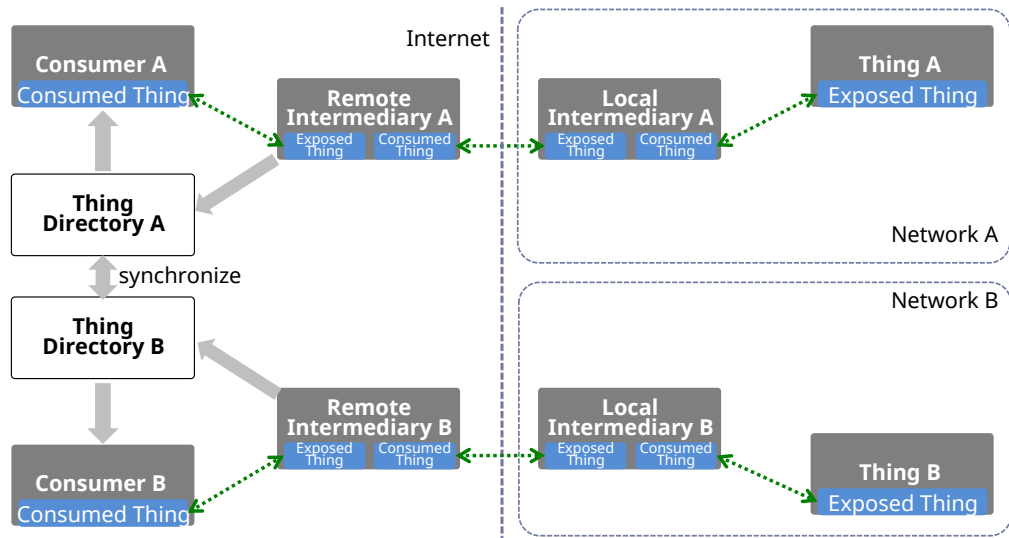


Figure 39 Multiple Cloud Connections Through Thing Description Directory Synchronization

#### § 9.2.5.2 Connection Through Proxy Synchronization

In [Figure 40](#), two remote [Intermediaries](#) synchronize device information. When a shadow of [Thing B](#) is created in remote [Intermediary B](#), the TD of the shadow is simultaneously synchronized into remote [Intermediary A](#). Remote [Intermediary A](#) in turn creates its own shadow of [Thing B](#), and registers the TD with [Thing Description Directory A](#). With this mechanism, synchronization between [Thing Description Directory](#) services is not necessary.

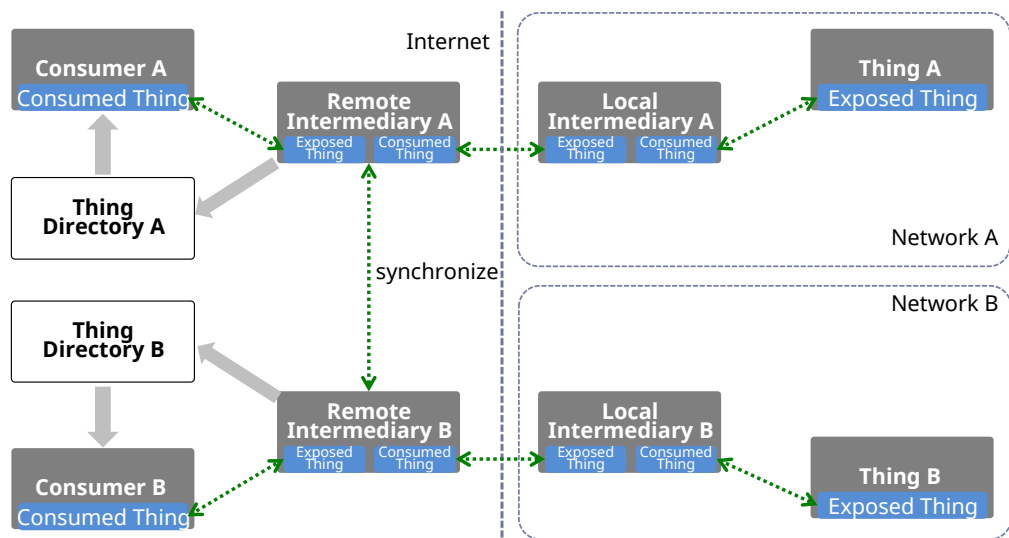


Figure 40 Multiple Cloud Connections Through Intermediary Synchronization

## § 10. Security Considerations

Security is a cross-cutting issue that needs to be considered in all [WoT building blocks](#) and WoT implementations. This chapter summarizes some general issues and guidelines to help preserve the security of concrete WoT implementations. However, these are only general guidelines and an abstract architecture such as described in this document cannot, itself, guarantee security. Instead the details of a concrete implementation need to be considered. For a more detailed and complete analysis of security (and privacy) issues, see the *WoT Security and Privacy Guidelines* document [[WOT-SECURITY](#)].

Overall, the goal of the WoT is to describe the existing access mechanisms and properties of IoT devices and services, including security. In general, ~~W3C~~ WoT is designed to describe what exists rather than to prescribe what to implement. A description of an existing system should accurately describe that system, even if it has less than ideal security behavior. A clear understanding of the security vulnerabilities of a system supports security mitigations—although of course such data need not be distributed to those who might maliciously exploit it.

However, especially for new systems, the WoT architecture should *enable* the use of best practices. In general, the WoT security architecture must support the goals and mechanisms of the IoT protocols and systems it connects to. These systems vary in their security requirements and risk tolerance, so security mechanisms will also vary based on these factors.

Security is especially important in the IoT domain since IoT devices need to operate autonomously and, in many cases, may be in control of safety-critical systems. IoT devices are subject to different and in some cases higher risks than IT systems. It is also important to protect IoT systems so that they cannot be used to launch attacks on other computer systems.

In general, security cannot be guaranteed. It is not possible for the WoT to turn an insecure system into a secure one. However, the WoT architecture needs to do no harm: it should support security at least as well as the systems it describes and supports.

### § 10.1 WoT Thing Description Risks

The metadata contained in a [WoT Thing Description](#) (TD) is potentially sensitive. As a best practice, TDs should be used together with appropriate integrity protection mechanisms and access control policies, with the goal of providing it only to authorized users.

Please refer to the Privacy Considerations sections of the WoT Thing Description specification for additional details and discussion of these points.

### § 10.1.1 Thing Description Private Security Data Risk

TDs are designed to carry only Public Security Metadata. The built-in TD security schemes defined in the TD specification intentionally do not support the encoding of Private Security Data.

However, there is a risk that other fields such as human-readable descriptions might be misused to encode this information, or new security schemes might be defined and deployed via the extension mechanism that encode such information.

#### **Mitigation:**

There *SHOULD* be a strict separation of Public Security Metadata and Private Security Data. Authentication and authorization *SHOULD* be established based on separately managed Private Security Data. Producers of TDs *MUST* ensure that no Private Security Data is included in TDs.

### § 10.1.2 Thing Description Communication Metadata Risk

Without best-practice configuration of security mechanisms, communication with IoT devices is at greater risk of being intercepted or compromised.

#### **Mitigation:**

Configure any communication security metadata for an IoT Platform used with WoT following best practices for that IoT Platform. Whenever possible, TD creators *SHOULD* use the vetted communication metadata provided in the WoT Binding Templates. When generating TDs for an IoT ecosystem not covered by the WoT Binding Templates, TD creators *SHOULD* ensure that all the security requirements of the IoT Platform are satisfied.

## § 10.2 WoT Scripting API Risks

The WoT Runtime implementation and the WoT Scripting API should have mechanisms to prevent malicious access to the system and isolate scripts in multi-tenant Servients. More specifically the WoT Runtime implementation when used with the WoT Scripting API should consider the following security and privacy risks and implement the recommended mitigations.

In general, these risks and mitigations should also be applied to any system that supports programmable behavior for WoT systems, not just the WoT Scripting API.

## § 10.2.1 Cross-Script Security Risk

In basic WoT setups, all scripts running inside the WoT Runtime are considered trusted, distributed by the manufacturer, and therefore there is no strong need to isolate script instances from each other. However, depending on device capabilities, deployment use case scenarios, and risk level it might be desirable to do so. For example, if one script handles sensitive data and is well-audited, it might be desirable to separate it from the rest of the script instances to minimize the risk of data exposure in case some other script inside the same system gets compromised during runtime. Another example is mutual co-existence of different tenants on a single WoT device. In this case each WoT runtime instance will be hosting a different tenant, and preventing tenants from accessing each other's data without authorization will generally be needed.

### **Mitigation:**

The WoT Runtime *SHOULD* perform isolation of script instances and their data from each other in cases when scripts handle sensitive data. Similarly, the WoT Runtime implementation *SHOULD* perform isolation of WoT Runtime instances and their data from each other if a WoT device has more than one tenant. In practice, isolation of scripts and runtime instances from each other can be accomplished by running all instances in a "sandboxed" environment that controls its access to the rest of the environment. For more information see Sections [WoT Servient Single-Tenant](#) and [WoT Servient Multi-Tenant](#) of the *WoT Security and Privacy Guidelines* specification [[WOT-SECURITY](#)].

## § 10.2.2 Physical Device Direct Access Risk

In case a script is compromised or malfunctions the underlying physical device (and potentially surrounded environment) can be damaged if a script can use directly exposed native device interfaces. If such interfaces lack safety checks on their inputs, they might bring the underlying physical device (or environment) to an unsafe state.

### **Mitigation:**

The WoT Runtime *SHOULD NOT* directly expose native device interfaces to the script developers. A WoT Runtime implementation *SHOULD* provide a hardware abstraction layer for accessing the native device interfaces. Hardware abstraction layers *SHOULD* refuse to execute commands that might put the device (or environment) to an unsafe state. Additionally, in order to reduce the damage to a physical WoT device in cases a script gets compromised, it is important to minimize the number of interfaces that are exposed or accessible to a particular script based on its functionality.



### § 10.3.1 Provisioning and Update Security Risk

If the WoT Runtime implementation supports post-manufacturing provisioning or updates of itself, scripts, or any related data (including security credentials), it can be a major attack vector. An attacker can try to modify any above described element during the update or provisioning process or simply provision attacker's code and data directly.

**Mitigation:**

Post-manufacturing provisioning or update of scripts, the WoT Runtime itself or any related data *SHOULD* be done in a secure fashion. A set of recommendations for secure update and post-manufacturing provisioning can be found in the *WoT Security and Privacy Guidelines* specification [[WOT-SECURITY](#)].

### § 10.3.2 Security Credentials Storage Risk

Typically the WoT Runtime needs to store the security credentials that are provisioned to a WoT device to operate in a network. If an attacker can compromise the confidentiality or integrity of these credentials, then it can obtain access to assets, impersonate other WoT Things, devices, or services, or launch Denial-Of-Service (DoS) attacks.

**Mitigation:**

The WoT Runtime *SHOULD* securely store any provisioned security credentials, guaranteeing their integrity and confidentiality. In case there are more than one tenant on a single WoT-enabled device, a WoT Runtime implementation *SHOULD* isolate each tenant's provisioned security credentials from other tenants. In order to minimize a risk that provisioned security credentials get compromised, the WoT Runtime implementation *SHOULD NOT* expose any API for scripts to query provisioned security credentials. Such credentials (or even better, abstract operations that use them but do not expose them) *SHOULD* only be accessible to the underlying protocol implementation that uses them.

## § 10.4 Trusted Environment Risks

In section 5. [Common Deployment Patterns](#) several usage scenarios are presented that include the concept of a Trusted Environment and a security boundary. Entities that are members of a Trusted Environment all share access to a common set of resources (such as a local network) and are

implicitly granted certain access rights to each other. A common example would be a WiFi LAN in the home where access to the WEP password allows devices to communicate with each other without any further access controls. Allowing implicit access rights like this and using a single shared secret for a large number of entities means that a single malicious actor with access to the Trusted Environment can cause significant damage.

One common IoT situation is the use of an HTTP/HTML browser to access locally hosted web services in a home environment. Such locally hosted web services may not have a publicly visible URL and so are not able to participate in the CA system expected by browsers to enable use of HTTP/TLS (HTTPS). It is often the case in this situation that plain HTTP is used and the only security protecting the communication is network encryption, for example WEP, which is relatively weak.

**Mitigation:**

Trust relationships *SHOULD* be as restricted as possible, ideally pairwise and limited to precisely the access required. As noted, in some situations this is difficult to manage and implicit access such as described is used, and may even be assumed by certain entities, such as browsers. In the case of implicit access control via access to a common network a segmented network *SHOULD* be used. For example, in the home environment, a separate WiFi network can be used for IoT devices, and routers often provide a "guest" network that can be used for this purpose. In commercial and industrial environments, explicit installation of pre-shared keys *SHOULD* be used to allow browsers to access local services while using TLS. Using a single key for a large number of services is equivalent to the "implicit access" situation above. However, once TLS is enabled other security mechanisms that assume secure transport can be used to provide fine-grained access control.

## § 10.5 Secure Transport

As noted in [10.4 Trusted Environment Risks](#), there are situations in which secure transports, such as TLS, are not feasible or are difficult to set up, such as on a local LAN within a home. Unfortunately, generally access control mechanisms in HTTP are designed to be used with secure transport and can be easily bypassed without it. In particular, it is relatively easy to capture passwords and tokens from unencrypted protocol interactions which can be intercepted by a third party. Also, man-in-the-middle attacks can be easily implemented without TLS providing server authentication.

**Mitigation:**

Because of the practical difficulties in setting up secure transport in all situations, we cannot make a blanket assertion that it is always required. Instead we provide a set of requirements for different use cases:

**Public Networks:**

When a Thing is made available on the public internet so it can be accessed by anyone, from anywhere, then it *MUST* be protected by secure transport such as TLS or DTLS. A Thing available on the public internet will have a public URL and so normal CA mechanisms to provide certificates are available. This should be done even if there are no access controls on the endpoint, for example when providing a publicly accessible service, because secure transport also protects the privacy of requesters (for example, the content of queries).

**Private Networks:**

When a Thing is made available on a private network then it *SHOULD* be protected by secure transport such as TLS or DTLS. The risk increases with the sensitivity of the data being protected, the potential for damage, and with the number of people given access to the private network. Secure transport is a higher priority in higher risk situations such as a factory network; in such cases, it is also more practical to install pre-shared keys. In low-risk situations, the following is permissible: Private networks such as a LAN, protected by a firewall, *MAY* use the [Trusted Environment](#) approach of depending on network security only. This is not generally recommended but may be necessary for practical reasons. Please see the referenced security consideration for additional risks and mitigations with this approach.

**Brownfield Devices:**

WoT is descriptive and an accurate description of existing "brownfield" devices may reveal they are not secure. For example, a specific device may use HTTP without TLS, which is not secure even if it uses access controls such as a password. It is often not possible to upgrade such devices to support stronger security. In these cases, the above two mitigations still apply. If exposed on a private network, the [Trusted Environment](#) approach should be used, with access limited as much as possible. If it is desired to expose such a device on the public internet, additional steps should be taken to enhance its security. In particular, a proxy can be used to ensure that communications over the public network use secure transport. A proxy can also be used to provide access controls.

When secure transport over TCP is appropriate, then at least TLS 1.3 [[RFC8446](#)] *SHOULD* be used. If TLS 1.3 cannot be used for compatibility reasons but secure transport over TCP is appropriate, TLS 1.2 [[RFC5246](#)] *MAY* be used. When secure transport over UDP is appropriate, then at least DTLS 1.3 [[RFC9147](#)] *SHOULD* be used. If DTLS 1.3 cannot be used for compatibility reasons but secure transport over UDP is appropriate, then DTLS 1.2 [[RFC6347](#)] *MAY* be used. Versions of DTLS or TLS earlier than 1.2 *MUST NOT* be used for new development. Existing Things using earlier versions of TLS or DTLS can be described by WoT metadata (e.g. Thing Descriptions) but should be considered insecure.

Additional considerations apply if a Thing can reveal [Personally Identifiable Information](#) (PII). See [11.2 Access to Personally Identifiable Information](#).

## § 11. Privacy Considerations

Privacy is a cross-cutting issue that needs to be considered in all [WoT building blocks](#) and WoT implementations. This chapter summarizes some general issues and guidelines to help preserve the privacy of concrete WoT implementations. However, these are only general guidelines and an abstract architecture such as described in this document cannot, itself, guarantee privacy. Instead the details of a concrete implementation need to be considered. For a more detailed and complete analysis of privacy (and security) issues, see the *WoT Security and Privacy Guidelines* specification [[WOT-SECURITY](#)].

### § 11.1 WoT Thing Description Risks

The metadata contained in a [WoT Thing Description](#) (TD) is potentially sensitive, and even if it does not explicitly contain [Personally Identifiable Information](#) (PII) it may be possible to infer PII from it. As a best practice, TDs should be used together with appropriate integrity protection mechanisms and access control policies, with the goal of providing it only to authorized users. In general, many of the [Security Considerations](#) discussed in the previous section can also be seen as privacy risks, when they relate the undesired an unauthorized disclosure of information.

Please refer to the Privacy Considerations sections of the WoT Thing Description specification for additional details and discussion of these points.

#### § 11.1.1 Thing Description Personally Identifiable Information Risk

Thing descriptions can potentially contain [Personally Identifiable Information](#) (PII) of various types. Even if it is not explicit, a TD and its association with an identifiable person can be used to infer information about that person. For example, the association of fingerprintable TDs exposed by mobile devices whose location can be determined can be a tracking risk. Even if a particular device instance cannot be identified, the type of device represented by a TD, when associated with a person, may constitute personal information. For example, a medical device may be used to infer that the user has a medical condition.

Generally, [Personally Identifiable Information](#) in a TD should be limited as much as possible. In some cases, however, it cannot be avoided. The potential presence of both direct and inferencable PII in a TD means that TD as a whole should be treated like other forms of PII. They should be stored and transmitted in a secure fashion, should only be provided to authorized users, should only

be cached for limited times, should be deleted upon request, should only be used for the purpose for which they were provided with user consent, and they should otherwise satisfy all requirements (including any legal requirements) for the use of PII.

**Mitigation:**

Storage of explicit PII in TDs *SHOULD* be minimized as much as possible. Even without explicit PII in TDs, a tracking and identification privacy risk may exist. TDs that can be associated with a person *SHOULD* generally be treated as if they contained PII and subject to the same management policies as other PII, even if they do not explicitly contain it.

Distribution mechanisms for TDs *SHOULD* ensure they are only provided to authorized Consumers. Note that the WoT Discovery mechanism is designed to address these specific issues, as long as it is used with authentication and access controls on exploration services. As a general matter of policy, unnecessary information should not be exposed in TDs whenever possible. For example, explicit type and instance identifying information in TDs should only be included if it is needed by the use case. Even if required by the use case, to minimize tracking risks, distributed and limited-scope identifiers should be used whenever possible rather than globally unique identifiers. Other forms of information, such as human-readable descriptions, may also be omitted in some use cases to reduce fingerprinting risks.

## § 11.2 Access to Personally Identifiable Information

In addition to the risks of revealing Personally Identifiable Information (PII) through metadata discussed in 11.1.1 Thing Description Personally Identifiable Information Risk, the data returned by Things can itself be sensitive. For example, a Thing could be monitoring the location or health of a specific person. Information associated with a person should be treated as PII even if it is not immediately obvious that it is sensitive, since it could be combined with other information to reveal sensitive information.

**Mitigation:**

Things returning data or metadata (such as TDs) associated with a person *SHOULD* use some form of access control. A special case of this is a Thing Description Directory, as described in [WOT-DISCOVERY], which is a Thing that returns Thing Descriptions as data. Such directory services are included in the above statement and should use access control if the TDs describe Things associated with identifiable people. In the case of services returning Thing Descriptions, the following also applies: Services returning Thing Descriptions with immutable IDs *SHOULD* use some form of access control. Following the principle that Thing Descriptions describing Things associated with specific persons should be treated as PII, even if they do not explicitly contain it, this implies that directories providing such TDs should use access control. Generally speaking, the only exceptions should be cases where access is controlled by another mechanism not described in the TD itself, such as a segmented network. Again it should also be noted that access controls are generally only effective when secure

transport is also used; see [10.5 Secure Transport](#). Use of access controls without secure transport, at best, only discourages casual access by unauthorized parties.

## § A. Recent Specification Changes

### § A.1 Changes from the WD published September 7, 2022

- Modify [Section 7, WoT Building Blocks](#) so that it is informative.
- Add terminology entries for Profile, System, and Directory.
- Updates to affiliation of Matthias Kovatsch (as a former editor).
- Add/update references to the WoT Use Cases and Requirements [[WOT-USE-CASES-REQUIREMENTS](#)] document.

### § A.2 Changes in WD published September 7, 2022 from the FPWD version

- Move Toru Kawaguchi to former editors.
- Update abstract to reflect use of prescription (as opposed to pure description) where necessary.
- New section under Terminology: Device Categories.
- Section removed: System Integration.
- Lifecycle section reworked; removal of separate "Information Lifecycle" section.
- Reference and discuss new WoT building block and document: WoT Discovery [[WOT-DISCOVERY](#)]
- Update discussion of WoT Profiles [[WOT-PROFILE](#)]
- Update discussion of WoT Binding Templates [[WOT-BINDING-TEMPLATES](#)]
- Rename "System Topologies (Horizontal)" to "Common Deployment Patterns".
- Rename "Application Domains (Vertical)" to "Application Domains".
- New or revised deployment patterns:
  - Telemetry
  - Orchestration
  - Virtual Things

- Revise Security Considerations and Privacy Considerations:
  - Reorganized to ensure that risks and mitigations are presented separately, and that normative content is only given under mitigations.
  - Added new section for Access to PII covering actual data access, not just metadata.
  - Added new section for Trusted Environments.
  - Added new section for Secure Transport. This section defines under what conditions (D)TLS should or must be used.

## § A.3 Changes in the FPWD from the 1.0 version of [wot-architecture]

- Make Security Considerations and Privacy Considerations normative.
- Add definitions for Connected Device (a.k.a. Device), Service, and Shadow. Update definitions for Virtual Thing and Digital Twin. Narrow definition of Virtual Thing to a Service that mediates with one of more other Things.
- Split Security and Privacy Considerations into separate chapters.
- New chapter: Lifecycle.
- New terminology: Thing Model.
- Chapter restructuring and renaming:
  - Application Domains (Verticals)
  - System Topologies (Horizontals)
  - System Integration
  - Abstract WoT System Architecture
- Various editors notes with placeholders for planned contributions.
- References to github repositories for use cases and requirements.
- Requirement chapter was moved to the *WoT Use Cases and Requirements* document [[WOT-USE-CASES-REQUIREMENTS](#)].

## § B. Acknowledgments

Special thanks to Matthias Kovatsch for co-editing and many significant contributions to the initial version of the WoT architecture specification. Special thanks to Michael McCool, Takuki Kamiya, Kazuyuki Ashimura, Sebastian Käbisch, Zoltan Kis, Ege Korkan, Michael Koster, Elena

Reshetova, Klaus Hartke, Ari Keränen, Kazuaki Nimura, and Philippe Le Hegaret for their contributions to this document.

Special thanks to Dr. Kazuyuki Ashimura from the W3C for the continuous help and support of the work of the WoT Architecture Task Force.

Many thanks to the W3C staff and all other active Participants of the W3C Web of Things Interest Group (WoT IG) and Working Group (WoT WG) for their support, technical input and suggestions that led to improvements to this document.

The WoT WG also would like to appreciate the pioneering efforts regarding the concept of "Web of Things" that started as an academic initiative in the form of publications such as [WOT-PIONEERS-1] [WOT-PIONEERS-2] [WOT-PIONEERS-3] [WOT-PIONEERS-4] and, starting in 2010, a yearly [International Workshop on the Web of Things](#).

Finally, special thanks to Joerg Heuer for leading the WoT IG for 2 years from its inception and guiding the group to come up with the concept of WoT building blocks including the Thing Description.

## § C. References

### § C.1 Normative references

#### [BCP47]

*Tags for Identifying Languages*. A. Phillips, Ed.; M. Davis, Ed.. IETF. September 2009. Best Current Practice. URL: <https://www.rfc-editor.org/rfc/rfc5646>

#### [JSON-LD]

*JSON-LD 1.0*. Manu Sporny; Gregg Kellogg; Markus Lanthaler. W3C. 3 November 2020. W3C Recommendation. URL: <https://www.w3.org/TR/json-ld/>

#### [RFC2046]

*Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types*. N. Freed; N. Borenstein. IETF. November 1996. Draft Standard. URL: <https://www.rfc-editor.org/rfc/rfc2046>

#### [RFC2119]

*Key words for use in RFCs to Indicate Requirement Levels*. S. Bradner. IETF. March 1997. Best Current Practice. URL: <https://www.rfc-editor.org/rfc/rfc2119>

#### [RFC3986]

*Uniform Resource Identifier (URI): Generic Syntax*. T. Berners-Lee; R. Fielding; L. Masinter. IETF. January 2005. Internet Standard. URL: <https://www.rfc-editor.org/rfc/rfc3986>



**[RFC3987]**

*Internationalized Resource Identifiers (IRIs)*. M. Duerst; M. Suignard. IETF. January 2005. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc3987>

**[RFC5246]**

*The Transport Layer Security (TLS) Protocol Version 1.2*. T. Dierks; E. Rescorla. IETF. August 2008. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc5246>

**[RFC6347]**

*Datagram Transport Layer Security Version 1.2*. E. Rescorla; N. Modadugu. IETF. January 2012. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc6347>

**[RFC8174]**

*Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words*. B. Leiba. IETF. May 2017. Best Current Practice. URL: <https://www.rfc-editor.org/rfc/rfc8174>

**[RFC8259]**

*The JavaScript Object Notation (JSON) Data Interchange Format*. T. Bray, Ed.. IETF. December 2017. Internet Standard. URL: <https://www.rfc-editor.org/rfc/rfc8259>

**[RFC8288]**

*Web Linking*. M. Nottingham. IETF. October 2017. Proposed Standard. URL: <https://httpwg.org/specs/rfc8288.html>

**[RFC8446]**

*The Transport Layer Security (TLS) Protocol Version 1.3*. E. Rescorla. IETF. August 2018. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc8446>

**[RFC9147]**

*The Datagram Transport Layer Security (DTLS) Protocol Version 1.3*. E. Rescorla; H. Tschofenig; N. Modadugu. IETF. April 2022. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc9147>

**[wot-architecture]**

*Web of Things (WoT) Architecture*. Matthias Kovatsch; Ryuichi Matsukura; Michael Lagally; Toru Kawaguchi; Kunihiro Tsumura; Kazuo Kajimoto. W3C. 9 April 2020. W3C Recommendation. URL: <https://www.w3.org/TR/wot-architecture/>

**[WOT-DISCOVERY]**

*Web of Things (WoT) Discovery*. Andrea Cimmino; Michael McCool; Farshid Tavakolizadeh; Kunihiro Tsumura. W3C. November 2020. URL: <https://www.w3.org/TR/wot-discovery/>

**[WOT-PROFILE]**

*Web of Things (WoT) Profile*. Michael Lagally; Michael McCool; Ryuichi Matsukura; Sebastian Käbis; Tomoaki Mizushima. W3C. 24 November 2020. W3C Working Draft. URL: <https://www.w3.org/TR/wot-profile/>

**[WOT-THING-DESCRIPTION]**

*Web of Things (WoT) Thing Description 1.1*. Sebastian Käbis; Takuki Kamiya; Michael McCool; Victor Charpenay. W3C. November 2020. URL: <https://www.w3.org/TR/wot-thing-description11/>

## [WOT-USE-CASES-REQUIREMENTS]

*Web of Things (WoT) Use Cases and Requirements*. Michael Lagally; Michael McCool; Ryuichi Matsukura; Tomoaki Mizushima. W3C. Oct 2020. Editor's Draft. URL: <https://www.w3.org/TR/wot-usecases/>

## § C.2 Informative references

### [CoRAL]

*The Constrained RESTful Application Language (CoRAL)*. Christian Amsüss; Thomas Fossati. IETF. March 2022. Internet-Draft. URL: <https://datatracker.ietf.org/doc/html/draft-ietf-core-coral/>

### [ECMAScript]

*ECMAScript Language Specification*. Ecma International. URL: <https://tc39.es/ecma262/multipage/>

### [HCI]

*The Encyclopedia of Human-Computer Interaction, 2nd Ed.* Interaction Design Foundation. 2013. URL: <https://www.interaction-design.org/literature/book/the-encyclopedia-of-human-computer-interaction-2nd-ed>

### [HTML]

*HTML Standard*. Anne van Kesteren; Domenic Denicola; Ian Hickson; Philip Jägenstedt; Simon Pieters. WHATWG. Living Standard. URL: <https://html.spec.whatwg.org/multipage/>

### [IANA-RELATIONS]

*Link Relations*. IANA. URL: <https://www.iana.org/assignments/link-relations/>

### [IEC-FOTF]

*Factory of the future*. IEC. October 2015. URL: <https://www.iec.ch/basecamp/factory-future/>

### [IOT-SCHEMA-ORG]

*Schema Extensions for IoT Community Group*. URL: <https://www.w3.org/community/iotschema/>

### [ISO-IEC-2382]

*Information technology — Vocabulary*. ISO. 2015. URL: <https://www.iso.org/obp/ui/#iso:std:iso-iec:2382:ed-1:v1:en>

### [ISO-IEC-27000]

*Information technology — Security techniques — Information security management systems — Overview and vocabulary*. ISO. 2018. URL: <https://www.iso.org/obp/ui/#iso:std:iso-iec:27000:ed-5:v1:en>

### [ISO-IEC-29100]

*Information technology — Security techniques — Privacy framework*. ISO. 2011. URL: <https://www.iso.org/obp/ui/#iso:std:iso-iec:29100:ed-1:v1:en>

### [JSON-LD11]

*JSON-LD 1.1*. Gregg Kellogg; Pierre-Antoine Champin; Dave Longley. W3C. 16 July 2020. W3C Recommendation. URL: <https://www.w3.org/TR/json-ld11/>

**[LINKED-DATA]**

*Linked Data Design Issues*. Tim Berners-Lee. W3C. 27 July 2006. W3C-Internal Document. URL: <https://www.w3.org/DesignIssues/LinkedData.html>

**[LWM2M]**

*Lightweight Machine to Machine Technical Specification: Core*. OMA SpecWorks. August 2018. Approved Version: 1.1. URL: [http://openmobilealliance.org/release/LightweightM2M/V1\\_1-20180710-A/OMA-TS-LightweightM2M\\_Core-V1\\_1-20180710-A.pdf](http://openmobilealliance.org/release/LightweightM2M/V1_1-20180710-A/OMA-TS-LightweightM2M_Core-V1_1-20180710-A.pdf)

**[MQTT]**

*MQTT Version 3.1.1 Plus Errata 01*. Andrew Banks; Rahul Gupta. OASIS Standard. December 2015. URL: <http://docs.oasis-open.org/mqtt/mqtt/v3.1.1/mqtt-v3.1.1.html>

**[NORMAN]**

*The Psychology of Everyday Things*. Donald A. Norman. Basic Books. 1988.

**[OCF]**

*OCF Core Specification*. Open Connectivity Foundation. April 2019. Version 2.0.2. URL: <https://openconnectivity.org/developer/specifications>

**[REST]**

*REST: Architectural Styles and the Design of Network-based Software Architectures*. Roy Thomas Fielding. University of California, Irvine. 2000. PhD thesis. URL: [https://www.ics.uci.edu/~fielding/pubs/dissertation/fielding\\_dissertation.pdf](https://www.ics.uci.edu/~fielding/pubs/dissertation/fielding_dissertation.pdf)

**[RFC6202]**

*Known Issues and Best Practices for the Use of Long Polling and Streaming in Bidirectional HTTP*. S. Loreto; P. Saint-Andre; S. Salsano; G. Wilkins. IETF. April 2011. Informational. URL: <https://www.rfc-editor.org/rfc/rfc6202>

**[RFC6690]**

*Constrained RESTful Environments (CoRE) Link Format*. Z. Shelby. IETF. August 2012. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc6690>

**[RFC7049]**

*Concise Binary Object Representation (CBOR)*. C. Bormann; P. Hoffman. IETF. October 2013. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc7049>

**[RFC7228]**

*Terminology for Constrained-Node Networks*. C. Bormann; M. Ersue; A. Keranen. IETF. May 2014. Informational. URL: <https://www.rfc-editor.org/rfc/rfc7228>

**[RFC7231]**

*Hypertext Transfer Protocol (HTTP/1.1): Semantics and Content*. R. Fielding, Ed.; J. Reschke, Ed.. IETF. June 2014. Proposed Standard. URL: <https://httpwg.org/specs/rfc7231.html>

**[RFC7252]**

*The Constrained Application Protocol (CoAP)*. Z. Shelby; K. Hartke; C. Bormann. IETF. June 2014. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc7252>

**[RFC7641]**

*Observing Resources in the Constrained Application Protocol (CoAP)*. K. Hartke. IETF. September 2015. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc7641>

**[SAREF]**

*Smart Appliances REference (SAREF) ontology*. ETSI. November 2015. URL: <https://sites.google.com/site/smartappliancesproject/ontologies/reference-ontology>

**[SOLID]**

*Solid Technical Reports*. W3C Solid CG. April 2021. URL: <https://solidproject.org/TR/>

**[VOCAB-SSN]**

*Semantic Sensor Network Ontology*. Armin Haller; Krzysztof Janowicz; Simon Cox; Danh Le Phuoc; Kerry Taylor; Maxime Lefrançois. W3C. 19 October 2017. W3C Recommendation. URL: <https://www.w3.org/TR/vocab-ssn/>

**[WOT-BINDING-TEMPLATES]**

*Web of Things (WoT) Binding Templates*. Michael Koster; Ege Korkan. W3C. 30 January 2020. W3C Working Group Note. URL: <https://www.w3.org/TR/wot-binding-templates/>

**[WOT-PIONEERS-1]**

*Mobile Service Interaction with the Web of Things*. E. Rukzio, M. Paolucci; M. Wagner, H. Berndt; J. Hamard; A. Schmidt. Proceedings of 13th International Conference on Telecommunications (ICT 2006), Funchal, Madeira island, Portugal. May 2006. URL: <https://pdfs.semanticscholar.org/3ee3/a2e8ce93fbf9ba14ad54e12adaeb1f3ca392.pdf>

**[WOT-PIONEERS-2]**

*Putting Things to REST*. Erik Wilde. UCB iSchool Report 2007-015, UC Berkeley, Berkeley, CA, USA. November 2007. URL: <http://dret.net/netdret/docs/wilde-irep07-015-restful-things.pdf>

**[WOT-PIONEERS-3]**

*Poster Abstract: Dyser – Towards a Real-Time Search Engine for the Web of Things*. Benedikt Ostermaier; B. Maryam Elahi; Kay Römer; Michael Fahrmaier; Wolfgang Kellerer. Proceedings of ACM SenSys 2008, Raleigh, NC, USA. November 2008. URL: <https://www.vs.inf.ethz.ch/publ/papers/ostermair-poster-2008.pdf>

**[WOT-PIONEERS-4]**

*A Resource Oriented Architecture for the Web of Things*. Dominique Guinard; Vlad Trifa; Erik Wilde. Proceedings of Internet of Things 2010 International Conference (IoT 2010). Tokyo, Japan. November 2010. URL: <https://ieeexplore.ieee.org/abstract/document/5678452>

**[WOT-SCRIPTING-API]**

*Web of Things (WoT) Scripting API*. Zoltan Kis; Daniel Peintner; Cristiano Aguzzi; Johannes Hund; Kazuaki Nimura. W3C. 24 November 2020. W3C Working Group Note. URL: <https://www.w3.org/TR/wot-scripting-api/>

**[WOT-SECURITY]**

*Web of Things (WoT) Security and Privacy Guidelines.* ; Michael McCool; Elena Reshetova.  
W3C. March 2019. URL: <https://www.w3.org/TR/wot-security/>

**[Y.4409-Y.2070]**

*ITU-T Rec. Y.4409/Y.2070 (01/2015) Requirements and architecture of the home energy management system and home network services.* ITU-T. January 2015. Recommendation.  
URL: <https://www.itu.int/rec/T-REC-Y.2070-201501-I>

